A factorised Pauli–Villars multigrid method for the Möbius domain-wall operator

Working note

October 18, 2026

The Möbius domain-wall operator has a spectrum that surrounds the origin, which defeats non-Hermitian Krylov solvers and most multigrid constructions. Kanamori, Chen and Matsufuru showed that preconditioning with the Pauli–Villars operator together with the site-diagonal part of the operator, $A' = (D_{\mathrm{PV}} C_{\mathrm{PV}}^{-1})^{\dagger} D C^{-1}$, moves the spectrum into the right half-plane and gives the smallest condition number among the operators they examined. This note describes a two-grid method for A'. The coarse-grid correction is factorised into one Galerkin solve for the quark operator and one for the adjoint Pauli–Villars operator, both with a nearest-neighbour coarse stencil, so that the outer system can be taken on the even sites without enlarging the coarse stencil; the prolongation is built from lifted approximate low modes for which the correction is exact on the vectors it is made from; the smoother is a fixed Chebyshev polynomial. We state the operator and its forms (with Neff's α), the correction, the coarse solvers, the setup, the smoother, the outer solver, and the questions that are open. The results (Sec. XI) are a mass scan on $24^3 \times 48$ against even-odd CGNE, the spectrum and pseudospectra of the preconditioned operator on 12^4 and 16^4 with the bound for GMRES they give, the behaviour of truncated GCR, and the sensitivity to the parameters; parts still being measured are marked.

Contents

I Introduction . . . 1
II Spectra of preconditioned Möbius operators . . . 3
III The operator and Pauli–Villars preconditioning . . . 4
IV Even-odd forms . . . 5
V The factorised coarse-grid correction and the prolongation . . . 6
VI The coarse solvers . . . 7
VII Setup . . . 8
VIII The Chebyshev smoother and the cycle . . . 9
IX The outer solver . . . 10
X Implementation in RutticeQCD . . . 10
XI Numerical results . . . 10
XII Open problems . . . 14

I. Introduction

Lattice computations of flavour physics and of the hadronic contributions to precision observables need a fermion discretisation whose chiral symmetry is broken only mildly at finite lattice spacing. The overlap operator [1, 2] has an exact lattice chiral symmetry. Domain-wall fermions [3, 4, 5] approximate it at a lower cost by adding a fifth dimension of extent L_s, in which modes of opposite chirality are bound to the two boundaries by a large negative Wilson mass $-M_5$; the residual chiral symmetry breaking is suppressed as L_s grows, and the effective four-dimensional operator is a rational approximation of the overlap operator. The Möbius generalisation [6, 7, 8] introduces two coefficients (b, c) and improves this approximation at fixed L_s; $c = 0$ is Shamir's operator.

The price is paid in the linear solver. The operator is neither Hermitian nor positive, one application costs L_s four-dimensional Wilson hops, and the standard solver is still the conjugate gradient method on the normal equation of the even-odd preconditioned operator, whose condition number is the square of that of the operator itself. For Wilson-type fermions this problem has been solved by multigrid: the low modes that slow down Krylov solvers are locally coherent [9], so a few of them, cut into blocks, span

the low-mode space, and adaptive aggregation multigrid makes the cost nearly independent of the quark mass [10, 11, 12].

Multigrid for domain-wall fermions. The domain-wall operator has a spectrum that encloses the origin, because of the large negative Wilson mass. This violates the half-plane condition of non-Hermitian Krylov methods, whose convergence is then of the order of the system size [13, 14]. Boyle [14] lists four multigrid approaches that have had a measure of success, and they differ in the operator that is coarsened. Cohen *et al.* [15] coarsened the squared unpreconditioned operator, with blocks covering the whole fifth dimension, and noted that its coarse operator reaches two hops. HDCG [16] accelerates the best standard solver, CG on the normal equation of the red-black (even-odd) operator, and coarsens that squared Schur complement; its coarse stencil is next-to-next-to-next-to-nearest neighbour, 81 points for blocks of at least 4^4. HDCR [17] works instead directly on the nearest-neighbour Hermitian indefinite operator $\gamma_5 R_5 D$ at $c = 0$, so that the coarse operator stays nearest neighbour, explicitly to make re-evaluation along an HMC trajectory affordable. Brower, Clark, Howarth and Weinberg [19] pointed out that the Pauli–Villars preconditioned operator $D_{\mathrm{PV}}^{-1} D$, not D, carries the physical four-dimensional degrees of freedom, solved the system $D_{\mathrm{PV}}^{\dagger} D x = D_{\mathrm{PV}}^{\dagger} b$ with coarse operators built from the projected Wilson kernel, and demonstrated it in the two-dimensional Schwinger model. Boyle and Yamaguchi [18] compared these approaches for Shamir fermions; HDCR could be set up and solved faster than a single red-black CG solve on large volumes. Two of their observations bear on the present method: among red-black solvers, CG on the Schur complement $M_{ee} - M_{eo} M_{oo}^{-1} M_{oe}$ needed far fewer iterations than on the left-preconditioned $1 - M_{ee}^{-1} M_{eo} M_{oo}^{-1} M_{oe}$; and combining a red-black smoother with the HDCR coarse correction was not effective, which they attributed to the eigenvectors of the red-black operator and of $\gamma_5 R_5 D$ not being aligned.

Recent work by Boyle and collaborators. The multiple right-hand-side HDCG of Ref. [14] works on the red-black Hermitian operator $(M_{ee} - M_{eo} M_{oo}^{-1} M_{oe})^{\dagger} (M_{ee} - M_{eo} M_{oo}^{-1} M_{oe})$ with a flexible ADEF2 two-level conjugate gradient, a coarse space spanned by blocked low modes obtained with Chebyshev filters, and a stationary coarse solver (deflation plus a Chebyshev polynomial). Treating several right-hand sides at once with block CG and batched matrix products makes the coarse space sub-dominant, and gains of up to a factor of twenty over red-black CG are reported at physical quark mass. In a Lattice 2026 talk, Boyle and Oare [20] named two limitations of this approach—the cost of forming the 81-point coarse operator, and that a single solve on a new gauge field, as in HMC, does not profit from many right-hand sides—and presented a multigrid that directly coarsens the two-hop operator $D_{\mathrm{PV}}^{\dagger} D$ with blocks that contain the whole fifth dimension, which gives a 33-point coarse stencil. They emphasised that $D_{\mathrm{PV}}^{\dagger} D$ is non-normal: its field of values can extend into the left half-plane although its spectrum does not, the Eisenstat–Elman–Schultz bound [21] for GCR is then lost, and they introduced a shift in the smoothers. The coarse solve dominated the cost; a three-level multiple right-hand-side version performs similarly to multiple right-hand-side HDCG, and an exact dense inverse on a further coarse level is planned.

Site-diagonal and Pauli–Villars preconditioning. Kanamori, Chen and Matsufuru [22] examined the spectra of several preconditioned Möbius operators. Combining the Pauli–Villars factor with the site-diagonal part C of the operator, they found that $A' = (D_{\mathrm{PV}} C_{\mathrm{PV}}^{-1})^{\dagger} D C^{-1}$ has eigenvalues with positive real part and the smallest condition number among the candidates, while replacing D_{PV} by a block-diagonal truncation brings back eigenvalues with negative real part. Independently, Neff's modification of the domain-wall matrix [23, 24] rescales the fifth dimension by a site-local factor Q and changes the conditioning of the even-odd system.

This note. We describe a two-grid method for A'. The outer system is the usual even-odd reduced one, as in HDCG, but the coarse-grid correction is *factorised*: the inverses of the quark operator and of the adjoint Pauli–Villars operator are replaced by two separate Galerkin inverses, each with a nearest-neighbour coarse stencil and an aggregation that is compatible with the five-dimensional chirality. The prolongation

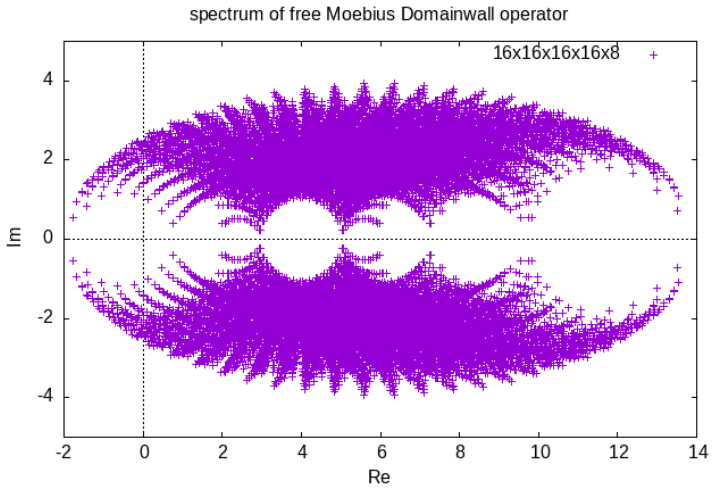


FIG. 1. Spectrum of a free Möbius domain-wall operator on $16^4 \times 8$ with quark mass 0.01, $b = 1.5$, $c = 0.5$ and domain-wall height $M_5 = 1$. Reproduced from Ref. [22], Fig. 1.

is built from approximate low modes lifted to the whole lattice in two ways, which makes the correction exact on these vectors for the right form of the operator. The smoother is a fixed Chebyshev polynomial and the setup combines Chebyshev filtering with a bootstrap. The note states the method; it does not assess its performance.

Section II recalls the spectra on which the construction rests. Sections III–IX describe the method: the operator and its Pauli–Villars preconditioning, its even-odd forms, the factorised correction and the prolongation, the coarse solvers, the setup, the smoother and the outer solver. Section XII lists the open problems. Statements that are our own derivations and have not been checked independently are marked **[derived here; not verified]**.

Conventions. Five-dimensional fields are $\psi_s(x)$, $s = 0, \ldots, L_s - 1$; $P_+, P_- = \frac{1}{2}(1 \pm \gamma_5)$. R reverses the fifth dimension, $(R\psi)_s = \psi_{L_s-1-s}$, and $\Gamma_5 = \gamma_5 R$. Subscripts e and o denote blocks on the even and odd four-dimensional sites; X_{pq} maps parity q to parity p. The injection ι puts a field on the even sites onto the whole lattice (zero on the odd sites), and $\pi = \iota^\dagger$ takes the even sites.

II. Spectra of preconditioned Möbius operators

The location of the spectrum decides which iterative methods can work. A Krylov residual polynomial p with $p(0) = 1$ that is small on the spectrum exists at low degree only if the spectrum can be enclosed by a region, such as an ellipse, that excludes the origin. The same holds for the Chebyshev smoother of Sec. VIII, which is defined by an ellipse in the right half-plane. The spectra of Ref. [22], reproduced in Figs. 1–3, show what the preconditioners do. The lattices are given in the captions.

The unpreconditioned operator. The free spectrum (Fig. 1) is a band that reaches into $\mathrm{Re}\,\lambda < 0$. On a $2 + 1$-flavour configuration at the physical point (Fig. 2, top left) the real parts lie in $[-1.29, 13.2]$. Conjugate gradients and BiCGStab therefore do not apply to D itself, and the practical choice has been CG on a positive definite squared operator.

Site-diagonal preconditioning. With the site-diagonal part C, whose inverse is a constant matrix in the fifth dimension and spin (Sec. B), DC^{-1} has real parts in $[-0.22, 2.22]$ (middle left): the extent shrinks, the sign problem remains.

Pauli–Villars preconditioning. The operator $D_{\rm PV}^{-1}D$, with $D_{\rm PV} = D(m=1)$, is a rational approximation of the overlap operator [19]. Because $D_{\rm PV}^{-1}$ is as expensive as a solve with D, Ref. [19] used $D_{\rm PV}^\dagger$ instead. $D_{\rm PV}^\dagger D$ has positive real parts (top right), and its low modes lie close to the imaginary axis with a positive shift of the order of the quark mass, as expected of an approximate overlap operator; its high end, however, reaches about 187, which an overlap operator would not show. Replacing $D_{\rm PV}$ by its block-diagonal approximation $B_{\rm PV}$, which drops the hops between 4^4 blocks, brings back negative real parts (middle and bottom right).

Both together. The operator $A' = (D_{\rm PV}C_{\rm PV}^{-1})^\dagger D C^{-1}$ has real parts in $[2.67\times10^{-4},\ 5.22]$ (Fig. 3, left); its condition number, 1.96×10^4, is smaller than the 2.34×10^4 of $D_{\rm PV}^\dagger D$, and the smallest among the operators examined in Ref. [22]. With $D_{\rm PV}$ replaced by $B_{\rm PV}$ the real parts become negative again (right). A preconditioner that truncates the Pauli–Villars factor to blocks would therefore inherit eigenvalues with negative real part; in the method below the factor is kept exact on the fine lattice and is approximated only through its Galerkin coarse operator (Sec. V).

Two caveats. First, the figures show the full-lattice operator A'; the even-odd operators of Sec. IV are related to it but their spectra are not part of these figures. Second, positivity of the real parts of the eigenvalues is not positivity of the field of values. For a non-normal operator the convergence bound of GCR [21] requires the Hermitian part $(A+A^\dagger)/2$ to be positive definite, which can fail while every eigenvalue lies in the right half-plane; this is the mechanism discussed for $D_{\rm PV}^\dagger D$ in Ref. [20]. Whether it occurs for A' and its even-odd forms is open (Sec. XII).

III. The operator and Pauli–Villars preconditioning

A. The Möbius operator

Let K be the hopping term of the Wilson operator,

$$(K\psi)(x) = \sum_{\mu=1}^{4}\Big[(1-\gamma_\mu)U_\mu(x)\psi(x+\hat\mu) + (1+\gamma_\mu)U_\mu^\dagger(x-\hat\mu)\psi(x-\hat\mu)\Big], \tag{3.1}$$

and $W = (4-M_5) - \frac{1}{2}K$ the Wilson operator with mass $-M_5$ and hopping coefficient $\frac{1}{2}$. The hop in the fifth dimension is

$$(Y\psi)_s = -P_-\psi_{s+1} - P_+\psi_{s-1}, \qquad \psi_{L_s} \equiv -m\,\psi_0, \quad \psi_{-1} \equiv -m\,\psi_{L_s-1}, \tag{3.2}$$

where m is the quark mass, which enters only through the wrap-around. The Möbius operator is

$$D = 1 + Y + W\,(b - cY) = C - \tfrac{1}{2}K\,(b-cY), \qquad C = d_+ + d_-Y, \tag{3.3}$$

with $d_+ = 1 + b(4-M_5)$ and $d_- = 1 - c(4-M_5)$. K acts on each slice s and does not commute with P_+ and P_-, so the order in Eq. (3.3) matters. The Pauli–Villars operator is the same operator at unit mass, $D_{\rm PV} = D|_{m=1}$, $C_{\rm PV} = C|_{m=1}$.

Γ_5-Hermiticity. Since $\Gamma_5 K \Gamma_5 = K^\dagger$, $\Gamma_5 Y \Gamma_5 = Y^\dagger$ and $\Gamma_5 C \Gamma_5 = C^\dagger$,

$$\Gamma_5 D \Gamma_5 - D^\dagger = -\tfrac{1}{2}\big[K^\dagger,\, b - cY^\dagger\big] = \tfrac{c}{2}\big[K^\dagger, Y^\dagger\big]. \tag{3.4}$$

The commutator is non-zero because K contains γ_μ and Y contains $P_\pm$. Thus $\Gamma_5 D \Gamma_5 = D^\dagger$ holds for Shamir's operator and fails at $O(c)$ for $c \neq 0$; the Hermitian operator of the Möbius case is $\Gamma_5 D_-^{-1} D$ with $D_- = 1 - cW$ [8], which involves an inverse that is not local. This matters for the coarse solvers (Sec. VI).

B. The site-diagonal part and its inverse

C acts within a four-dimensional site, as an $L_s \times L_s$ matrix in the fifth dimension on each chirality. With $\psi^{\pm} = P_{\pm}\psi$, it is lower bidiagonal on ψ^{+} and upper bidiagonal on ψ^{-}, each with one corner element from the mass:

$$(C\psi)^{+}_{s} = \begin{cases} d_{+}\psi^{+}_{0} + m\, d_{-}\psi^{+}_{L_s-1}, & s = 0, \\ d_{+}\psi^{+}_{s} - d_{-}\psi^{+}_{s-1}, & s \geq 1, \end{cases} \tag{3.5}$$

and the same for ψ^{-} with $s \to L_s - 1 - s$. With $q = d_{-}/d_{+}$, the solution of $C^{+}x = r$ is a forward substitution $p_0 = 0$, $p_s = (r_s + d_{-}p_{s-1})/d_{+}$, followed by

$$x_0 = \frac{r_0 - m\, d_{-}\, p_{L_s-1}}{d_{+} + m\, d_{-}\, q^{L_s-1}}, \qquad x_s = p_s + q^{s}x_0; \tag{3.6}$$

for C^{-} the substitution runs from $s = L_s - 1$ downwards, and for $C^{\dagger}$ the roles of the two chiralities are exchanged. C^{-1} is thus exact and costs $O(L_s)$ per site. It does not depend on the gauge field [22].

C. The preconditioned operator

The operator of Ref. [22] is

$$A' = (D_{\mathrm{PV}}C_{\mathrm{PV}}^{-1})^{\dagger}\, D\, C^{-1}, \qquad A'y = (D_{\mathrm{PV}}C_{\mathrm{PV}}^{-1})^{\dagger}\,\eta, \quad \psi = C^{-1}y, \tag{3.7}$$

for $D\psi = \eta$. For $m \to 1$ it becomes the Hermitian positive operator $(D_{\mathrm{PV}}C_{\mathrm{PV}}^{-1})^{\dagger}(D_{\mathrm{PV}}C_{\mathrm{PV}}^{-1})$.

D. Neff's α

Neff's form of the domain-wall matrix [23, 24] is $D^{(\alpha)} = DQ$ with a site-local Q, diagonal in the fifth dimension,

$$Q_0 = P_{-} + \alpha P_{+}, \qquad Q_s = \alpha \;\; (0 < s < L_s - 1), \qquad Q_{L_s-1} = P_{+} + \alpha P_{-}. \tag{3.8}$$

Q is Hermitian, Q^{-1} is Eq. (3.8) with $\alpha \to 1/\alpha$, and Q leaves $P_{-}\psi_0$ and $P_{+}\psi_{L_s-1}$ unchanged. The site-diagonal part of DQ is CQ, with $(CQ)^{-1} = Q^{-1}C^{-1}$ and $(CQ)^{\dagger} = QC^{\dagger}$. The solution of $D\psi = \eta$ is $\psi = Q\psi^{(\alpha)}$ with $D^{(\alpha)}\psi^{(\alpha)} = \eta$. The factor may be applied to the quark operator alone or to the Pauli–Villars operator as well; we write $D_{\mathrm{PV}}^{(\alpha)} = D_{\mathrm{PV}}Q$ in the latter case.

IV. Even-odd forms

With $D = \begin{pmatrix} C_e & D_{eo} \\ D_{oe} & C_o \end{pmatrix}$, $D_{pq} = -\frac{1}{2}K_{pq}(b - cY)$, the Schur complement on the even sites is $S = C_e - D_{eo}C_o^{-1}D_{oe}$, and S_{PV} is that of D_{PV}; note $S^{-1} = \pi D^{-1}\iota$ and $(S_{\mathrm{PV}}^{\dagger})^{-1} = \pi(D_{\mathrm{PV}}^{\dagger})^{-1}\iota$. The site-diagonal factor can stand on either side of S, which gives the right and the left form,

$$\tilde{A}_R = (S_{\mathrm{PV}}C_{\mathrm{PV},e}^{-1})^{\dagger}\,(S\,C_e^{-1}), \qquad \tilde{A}_L = (C_{\mathrm{PV},e}^{-1}S_{\mathrm{PV}})^{\dagger}\,(C_e^{-1}S). \tag{4.1}$$

$\tilde{A}_R$ is the even-odd counterpart of A'; $C_e^{-1}S$ is the usual left preconditioned even-odd operator, which Boyle and Yamaguchi found inferior to S itself for CG on the normal equation [18]. With α, DQ has Schur complement SQ_e and diagonal block C_eQ_e, so $\tilde{A}_R$ is independent of α, whereas

$$\tilde{A}_L^{(\alpha)} = \begin{cases} Q_e\,(C_{\mathrm{PV},e}^{-1}S_{\mathrm{PV}})^{\dagger}\,Q_e^{-2}\,(C_e^{-1}S)\,Q_e, & \alpha \text{ in } D \text{ and } D_{\mathrm{PV}}, \\ (C_{\mathrm{PV},e}^{-1}S_{\mathrm{PV}})^{\dagger}\,Q_e^{-1}\,(C_e^{-1}S)\,Q_e, & \alpha \text{ in } D \text{ only}. \end{cases} \tag{4.2}$$

For $D\psi = \eta$, with $\eta'_e = \eta_e - D_{eo}C_o^{-1}\eta_o$, the right-hand side is $\rho_R = (S_{\mathrm{PV}}C_{\mathrm{PV},e}^{-1})^{\dagger}\eta'_e$ or $\rho_L = (C_{\mathrm{PV},e}^{-1}S_{\mathrm{PV}})^{\dagger}C_e^{-1}\eta'_e$ (operators with Q when $\alpha \neq 1$); the solution is $\psi_e^{(\alpha)} = C_e^{-1}y$ (right) or y (left), $\psi_o^{(\alpha)} = C_o^{-1}(\eta_o - D_{oe}\psi_e^{(\alpha)})$, $\psi = Q\psi^{(\alpha)}$.

TABLE 1. Site-local factors of the correction (5.2). They follow from $\tilde{A}^{-1}$ with $S^{-1} = \pi D^{-1}\iota$; with α the site-diagonal blocks are $C_e Q_e$ and $Q_e C^\dagger_{\mathrm{PV},e}$.

form	O_R	O_M	O_L
right, any α	$C^\dagger_{\mathrm{PV},e}$	1	C_e
left, α in D only	1	$C_e Q_e C^\dagger_{\mathrm{PV},e}$	Q_e^{-1}
left, α in D and D_{PV}	Q_e^{-1}	$C_e Q_e^2 C^\dagger_{\mathrm{PV},e}$	Q_e^{-1}

V. The factorised coarse-grid correction and the prolongation

A. Why factorise

Solving on the even sites halves the problem and improves the conditioning, but in the literature it has come at a price: the Schur complement is a two-hop operator, so HDCG, which coarsens the squared Schur complement, has an 81-point coarse stencil [16], while the methods with a more compact coarse operator coarsen something other than the Schur complement: $\gamma_5 R_5 D$ (HDCR [17]) or $D^\dagger_{\mathrm{PV}} D$ [19, 20]. The factorised correction avoids this choice: the outer operator is the even-odd one, while the coarse operators are those of D and D_{PV} themselves. By $S^{-1} = \pi D^{-1}\iota$, the inverse of the right form is

$$\tilde{A}_R^{-1} = C_e\, S^{-1}\, (S^\dagger_{\mathrm{PV}})^{-1}\, C^\dagger_{\mathrm{PV},e} = C_e\, \pi D^{-1}\iota\; \pi (D^\dagger_{\mathrm{PV}})^{-1}\iota\; C^\dagger_{\mathrm{PV},e}, \tag{5.1}$$

a product of two inverses of one-hop operators with site-local factors. The Galerkin operator of $\tilde{A}_R$ itself, a four-hop operator on the fine lattice, would have a much larger coarse stencil; the Galerkin operators of D and D_{PV} are nearest-neighbour, because D is nearest-neighbour in four dimensions and every aggregate contains the whole fifth dimension. We therefore replace each inverse by its Galerkin approximation $\mathcal{P}(\mathcal{P}^\dagger X \mathcal{P})^{-1}\mathcal{P}^\dagger$. For both forms and any α, the correction of a residual r on the even sites is

$$e = O_L\; \pi\mathcal{P}\, \mathcal{D}_{\mathrm{c}}^{-1}\mathcal{P}^\dagger\iota\;\; O_M\;\; \pi\mathcal{P}\left(\mathcal{P}^\dagger D^\dagger_{\mathrm{PV}}\mathcal{P}\right)^{-1}\mathcal{P}^\dagger\iota\;\; O_R\, r, \qquad \mathcal{D}_{\mathrm{c}} = \mathcal{P}^\dagger D \mathcal{P}, \tag{5.2}$$

with the site-local factors of Table 1. The coarse operators are always those of the plain operators D and D_{PV} (no α), so one setup serves both forms and every α.

B. Aggregation and prolongation

The fine lattice is divided into rectangular four-dimensional blocks; each block contains the whole fifth dimension and becomes one coarse site. Each test vector v is split into its Γ_5 eigencomponents $v_\pm = \frac{1}{2}(1 \pm \Gamma_5)v$, and within each block the $+$ parts and the $-$ parts are orthonormalised separately. With n_v test vectors a coarse site carries $2n_v$ components, the first n_v of positive and the last n_v of negative Γ_5-chirality. Writing $\Sigma = \mathrm{diag}(\mathbf{1}_{n_v}, -\mathbf{1}_{n_v})$ on each coarse site,

$$\Gamma_5 \mathcal{P} = \mathcal{P}\Sigma, \tag{5.3}$$

so that $\mathcal{P}^\dagger(\Gamma_5 X \Gamma_5)\mathcal{P} = \Sigma\, \mathcal{P}^\dagger X \mathcal{P}\, \Sigma$ for any fine operator X; in particular $\mathcal{P}^\dagger D^\dagger_{\mathrm{PV}}\mathcal{P} = \Sigma \mathcal{D}_{\mathrm{c}}^{\mathrm{PV}}\Sigma$ with $\mathcal{D}_{\mathrm{c}}^{\mathrm{PV}} = \mathcal{P}^\dagger D_{\mathrm{PV}}\mathcal{P}$ whenever $\Gamma_5 D_{\mathrm{PV}}\Gamma_5 = D^\dagger_{\mathrm{PV}}$, i.e. for $c = 0$. Because $v_\pm(L_s - 1 - s) = \pm\gamma_5 v_\pm(s)$, each part is fixed by half of the fifth dimension, so $\mathcal{P}$ need only be stored for $s < L_s/2$. On the even sites the transfer operators are $\mathcal{P}^\dagger\iota$ and $\pi\mathcal{P}$.

C. The basis $\{s, \hat{t}\}$

The prolongation is built from vectors y_j on the even sites that approximate low modes of $\tilde{A}_R$ (Sec. VII). Each y is lifted in two ways:

$$s = \left(C_e^{-1}y,\; -C_o^{-1}D_{oe}C_e^{-1}y\right), \qquad Ds = \iota\, S C_e^{-1} y, \tag{5.4}$$

$$\hat{t} = \left(a,\; -(C^\dagger_{\mathrm{PV},o})^{-1}(D_{\mathrm{PV},eo})^\dagger a\right), \quad a = \pi D s, \qquad D^\dagger_{\mathrm{PV}}\hat{t} = \iota\, S^\dagger_{\mathrm{PV}} a, \tag{5.5}$$

where the right-hand columns follow from the block form of D and of $D^\dagger_{\rm PV}$, whose odd-even block is $(D_{{\rm PV},eo})^\dagger$ and whose odd diagonal block is $C^\dagger_{{\rm PV},o}$. The images of s under D and of $\hat{t}$ under $D^\dagger_{\rm PV}$ thus vanish on the odd sites. The $2n$ vectors $\{s_j, \hat{t}_j\}$ are the test vectors, $n_v = 2n$.

Proposition 1 (Exactness, right form)**.** *Let $s, \hat{t} \in \operatorname{range}\mathcal{P}$, let $\mathcal{D}_{\rm c}$ and $\mathcal{P}^\dagger D^\dagger_{\rm PV}\mathcal{P}$ be invertible, and let both coarse systems in Eq.* (5.2) *be solved exactly with the exact adjoint $\mathcal{P}^\dagger D^\dagger_{\rm PV}\mathcal{P} = (\mathcal{D}^{\rm PV}_{\rm c})^\dagger$. Then the right-form correction of $r = \tilde{A}_R y$ is $e = y$.*

Proof. For any invertible $\mathcal{P}^\dagger X\mathcal{P}$, $\mathcal{P}(\mathcal{P}^\dagger X\mathcal{P})^{-1}\mathcal{P}^\dagger X$ is the identity on $\operatorname{range}\mathcal{P}$. Now $O_R r = C^\dagger_{{\rm PV},e}\tilde{A}_R y = S^\dagger_{\rm PV} S C_e^{-1} y = S^\dagger_{\rm PV} a$, and $\iota S^\dagger_{\rm PV} a = D^\dagger_{\rm PV}\hat{t}$ by Eq. (5.5). The first coarse solve returns $\pi\hat{t} = a = \pi D s$. With $O_M = 1$ the second has $\mathcal{P}^\dagger\iota\pi D s = \mathcal{P}^\dagger D s$ on the right, since Ds vanishes on the odd sites, and returns $\pi s = C_e^{-1} y$; $O_L = C_e$ gives y. □

Remark 1 (Conditions in practice). **[derived here; not verified]** The proposition needs (i) $s, \hat{t}$ in the range of the stored $\mathcal{P}$, which holds only up to the rounding of the reduced-precision storage of $\mathcal{P}$; (ii) the exact adjoint of the Pauli–Villars coarse factor—the Σ-adjoint $\Sigma\mathcal{D}^{\rm PV}_{\rm c}\Sigma$ of Sec. VI equals it only if $\Gamma_5 D_{\rm PV}\Gamma_5 = D^\dagger_{\rm PV}$, which by Eq. (3.4) fails for $c \neq 0$; (iii) exact coarse solves, whereas they are stopped at a loose tolerance. How much each departure costs has not been analysed.

Remark 2 (Left form). **[derived here; not verified]** For the left form with α in both operators, the same argument requires the quark lift s' of $Q_e y'$ and the Pauli–Villars lift $\hat{t}'$ of $u = C^{\dagger\,-1}_{{\rm PV},e} Q_e^{-2} C_e^{-1} S Q_e y'$ to lie in $\operatorname{range}\mathcal{P}$ (Table 1 gives $O_M u = S Q_e y'$). For $\alpha = 1$ and $y' = C_e^{-1} y$ one has $s' = s$, but $u = C^{\dagger\,-1}_{{\rm PV},e} C_e^{-1} a$ differs from a by a site-local matrix that mixes s and spin, so $\hat{t}'$ is in general not in the span of the right-form vectors. We therefore expect the correction not to be exact for the left form with the basis above, but this has not been verified.

D. The coarse operators

$\mathcal{D}_{\rm c} = \mathcal{P}^\dagger D\mathcal{P}$ and $\mathcal{D}^{\rm PV}_{\rm c} = \mathcal{P}^\dagger D_{\rm PV}\mathcal{P}$ are assembled from the hops of the fine operator between neighbouring blocks and from its block-internal part: on each coarse site a $2n_v \times 2n_v$ diagonal block and one link block per direction, nine blocks in all. The adjoint factor $\mathcal{P}^\dagger D^\dagger_{\rm PV}\mathcal{P}$ is not assembled; it is applied either exactly, as the conjugate transpose $(\mathcal{D}^{\rm PV}_{\rm c})^\dagger$ of the stored blocks, or approximately as $\Sigma\mathcal{D}^{\rm PV}_{\rm c}\Sigma$. When all extents of the coarse lattice are even, the coarse hops connect coarse sites of opposite parity and both coarse operators admit an even-odd decomposition of their own.

VI. The coarse solvers

Both coarse systems are solved approximately, to a loose relative residual or a maximum number of iterations, in reduced precision; reaching the tolerance is not required. They are solved in sequence: first the Pauli–Villars factor, whose solution enters the right-hand side of the quark factor.

A. The quark factor

Write $\mathcal{D}_{\rm c} = \begin{pmatrix} A_{ee} & B_{eo} \\ B_{oe} & A_{oo} \end{pmatrix}$ on the coarse lattice. The system $\mathcal{D}_{\rm c} c = g$ is solved through its Schur complement,

$$\left(A_{ee} - B_{eo}A_{oo}^{-1}B_{oe}\right) c_e = g_e - B_{eo}A_{oo}^{-1} g_o, \qquad c_o = A_{oo}^{-1}(g_o - B_{oe}c_e), \tag{6.1}$$

with BiCGStab or truncated flexible GCR, preconditioned by the exact inverses of the diagonal blocks A_{ee}^{-1}. $\mathcal{D}_{\rm c}$ is not Hermitian and $\Sigma\mathcal{D}_{\rm c}$ is Hermitian only for $c = 0$, so a non-Hermitian Krylov method is used.

B. The Pauli–Villars factor

The system is $(\mathcal{P}^\dagger D^\dagger_{\rm PV}\mathcal{P})\, w = g$. Two treatments are available.

Exact adjoint. Solve $(\mathcal{D}_{\mathrm{c}}^{\mathrm{PV}})^{\dagger} w = g$ with a non-Hermitian Krylov method, preconditioned by the inverses of the adjoint diagonal blocks, $(A_{xx}^{\dagger})^{-1} = (A_{xx}^{-1})^{\dagger}$.

Σ-Hermitian form. Replace the adjoint by $\Sigma\mathcal{D}_{\mathrm{c}}^{\mathrm{PV}}\Sigma$ and solve the equivalent system

$$\mathcal{H}_{\mathrm{c}}\, w = \Sigma g, \qquad \mathcal{H}_{\mathrm{c}} = \mathcal{D}_{\mathrm{c}}^{\mathrm{PV}}\Sigma, \tag{6.2}$$

with MINRES. $\mathcal{H}_{\mathrm{c}}$ is Hermitian (indefinite) exactly when $(\mathcal{D}_{\mathrm{c}}^{\mathrm{PV}})^{\dagger} = \Sigma\mathcal{D}_{\mathrm{c}}^{\mathrm{PV}}\Sigma$, i.e. for $c = 0$ by Eqs. (3.4) and (5.3); for $c \neq 0$ MINRES is applied to a non-Hermitian matrix and has no guarantee. The preconditioner must be Hermitian positive definite; we use the absolute value of the diagonal blocks [25], $|\mathcal{H}_{xx}|^{-1} = V|\Lambda|^{-1}V^{\dagger}$ with $\frac{1}{2}(\mathcal{H}_{xx} + \mathcal{H}_{xx}^{\dagger}) = V\Lambda V^{\dagger}$ on each coarse site. On a coarse lattice with even extents, the Schur complement of $\mathcal{H}_{\mathrm{c}}$ on the even sites is $S_{\mathrm{c}}\Sigma_e$, with S_{c} the Schur complement of $\mathcal{D}_{\mathrm{c}}^{\mathrm{PV}}$; one solves for $u_e = \Sigma w_e$ on the even sites, completes u on the odd sites as in Eq. (6.1), and sets $w = \Sigma u$.

C. Levels below the coarse lattice

Both factors may be preconditioned by further levels. The coarse sites are grouped into super-blocks; on each super-block, local vectors are formed from the coarse test vectors $\mathcal{P}^{\dagger}s_j$ and $\mathcal{P}^{\dagger}\hat{t}_j$, split by Σ and orthonormalised (Gram–Schmidt applied twice); the level operators are obtained by probing the nearest-neighbour stencil; the last level is inverted densely. With Q_ℓ the prolongation into level $\ell - 1$ and B_ℓ the diagonal blocks of level ℓ:

Quark factor. A V-cycle with the coarse correction followed by one block-Jacobi sweep on the updated residual, without pre-smoothing,

$$z = Q_{\ell+1}M_{\ell+1}Q_{\ell+1}^{\dagger}r, \qquad M_\ell r = z + B_\ell^{-1}(r - \mathcal{D}_\ell z); \tag{6.3}$$

on the coarse lattice itself the same multiplicative combination is used, and with the even-odd coarse form its even part preconditions the Schur system. The level test vectors can optionally be refined by a few GCR steps preconditioned by the levels already built.

Pauli–Villars factor (Σ-Hermitian form only). The additive Hermitian positive preconditioner

$$M_\ell = |B_\ell|^{-1} + Q_{\ell+1}M_{\ell+1}Q_{\ell+1}^{\dagger}, \qquad M_{\ell_{\max}} = |\mathcal{H}_{\ell_{\max}}|^{-1}, \tag{6.4}$$

in the spirit of absolute-value preconditioning [25].

VII. Setup

The setup produces the vectors $y_1, \ldots, y_n$ on the even sites, then the basis $\{s_j, \hat{t}_j\}$, $\mathcal{P}$, and the coarse operators. It always uses the right form; the same setup serves the left forms and every α (Sec. V).

A. Chebyshev-filtered subspace iteration

The vectors approximate the right singular vectors of the smallest singular values of $\tilde{A}_R$, i.e. eigenvectors of the Hermitian positive operator $H_n = \tilde{A}_R^{\dagger}\tilde{A}_R$. With the Chebyshev polynomial T_d and $\ell(x) = \big(2x - (\lambda_{\mathrm{hi}} + \lambda_{\mathrm{lo}})\big)/(\lambda_{\mathrm{hi}} - \lambda_{\mathrm{lo}})$, $|T_d(\ell(x))| \leq 1$ on $[\lambda_{\mathrm{lo}}, \lambda_{\mathrm{hi}}]$ while components below λ_{lo} grow; $T_d(\ell(H_n))$ is applied by the three-term recurrence.

Algorithm 1. LOWMODES: Chebyshev-filtered subspace iteration on $H_n = \tilde{A}_R^\dagger \tilde{A}_R$.

$X \leftarrow n$ random vectors on the even sites

$\lambda_{\rm hi} \leftarrow$ a bound slightly above $\lambda_{\max}(H_n)$ (power iteration); $\lambda_{\rm lo} \leftarrow$ a small fraction of $\lambda_{\rm hi}$

for each round **do**

$d \leftarrow$ a degree proportional to $\sqrt{\lambda_{\rm hi}/\lambda_{\rm lo}}$, within fixed bounds

$X \leftarrow T_d(\ell(H_n))\, X$; orthonormalise X through its Gram matrix

$G \leftarrow X^\dagger H_n X = V\Theta V^\dagger$; $X \leftarrow XV$ ▷ Rayleigh–Ritz

$\lambda_{\rm lo} \leftarrow \theta_{\max}$ ▷ largest Ritz value

end for

return X, Θ

B. Bootstrap

Each bootstrap round builds the basis and the two-grid cycle of Sec. VIII from the current y, replaces every y_j by the result of a fixed small number of FGMRES steps on $\tilde{A}_R x = y_j$ from $x = 0$ preconditioned by that cycle, and orthonormalises the new y through their Gram matrix ($y \leftarrow yV\Lambda^{-1/2}$ with $y^\dagger y = V\Lambda V^\dagger$). The cycle damps the components it already handles, so the new vectors are enriched in what it does not.

Algorithm 2. Setup of the two-grid method.

$y \leftarrow$ LOWMODES(H_n)

for each bootstrap round **do**

build $\mathcal{P}$ from $\{s_j, \hat{t}_j\}$ of y and the cycle $\mathcal{M}$ on it

for $j = 1, \ldots, n$ **do**

$y_j \leftarrow$ a few FGMRES steps on $\tilde{A}_R x = y_j$, preconditioned by $\mathcal{M}$

end for

orthonormalise y

end for

compute $s_j, \hat{t}_j$ by Eqs. (5.4)–(5.5); split by Γ_5; orthonormalise per block $\to \mathcal{P}$

assemble $\mathcal{D}_{\rm c}$, $\mathcal{D}_{\rm c}^{\rm PV}$, the inverses of their diagonal blocks, their even-odd forms

optionally, build the levels below from $\mathcal{P}^\dagger s_j$, $\mathcal{P}^\dagger \hat{t}_j$

VIII. The Chebyshev smoother and the cycle

A. Smoother

The smoother is ν steps of Chebyshev iteration [26, 27] for $\tilde{A}x = b$ from $x = 0$, i.e. a fixed polynomial $x = p(\tilde{A})b$ of degree $\nu - 1$. It is defined by an ellipse with real foci $[\lambda_{\rm lo}, \lambda_{\rm hi}]$: $\lambda_{\rm hi}$ slightly above an estimate of the spectral radius of $\tilde{A}$ from a power iteration, and $\lambda_{\rm lo} = \lambda_{\rm hi}/\kappa_s$ for a chosen ratio κ_s. With $\theta = \frac{1}{2}(\lambda_{\rm hi} + \lambda_{\rm lo})$, $\delta = \frac{1}{2}(\lambda_{\rm hi} - \lambda_{\rm lo})$, $\sigma = \theta/\delta$, $\varrho_0 = 1/\sigma$, $r_0 = b$ and $d_0 = r_0/\theta$, step k is

$$x \leftarrow x + d_k, \quad r_{k+1} = r_k - \tilde{A} d_k, \quad \varrho_{k+1} = \frac{1}{2\sigma - \varrho_k}, \quad d_{k+1} = \varrho_{k+1}\varrho_k\, d_k + \frac{2\varrho_{k+1}}{\delta}\, r_{k+1}, \tag{8.1}$$

the last step only updating x; ν steps cost $\nu - 1$ applications of $\tilde{A}$. Only a real focal interval is used.

B. Cycle

Algorithm 3. Two-grid cycle $y = \mathcal{M}b$ for $\tilde{A}$.

1: $y \leftarrow p(\tilde{A})\, b$ ▷ pre-smoothing
2: $r \leftarrow b - \tilde{A}y$
3: $y \leftarrow y + e(r)$ ▷ correction (5.2)
4: $r \leftarrow b - \tilde{A}y$
5: $y \leftarrow y + p(\tilde{A})\, r$ ▷ post-smoothing

One cycle costs 2ν applications of $\tilde{A}$ (two smoothings and two residuals) plus the two coarse solves. Because the coarse solves are Krylov iterations stopped by a residual test, $\mathcal{M}$ is not a fixed linear map.

IX. The outer solver

The outer iteration solves $\tilde{A}y = \rho$ with the cycle as right preconditioner, and ψ is reconstructed as in Sec. IV. Since $\mathcal{M}$ varies from one application to the next, flexible methods are the ones with a theoretical footing: restarted FGMRES and truncated flexible GCR, which keep the preconditioned directions; BiCGStab with the cycle as preconditioner is also available but has no flexible counterpart of its convergence theory. Block variants treat several right-hand sides together. Convergence is monitored on the even-odd system; the residual of the original equation, $\|D\psi - \eta\|/\|\eta\|$, is evaluated with the plain D after reconstruction. The reference solver is CG on the normal equation of the even-odd operator, which is robust but squares the condition number.

X. Implementation in RutticeQCD

The method is implemented in RutticeQCD (Rust; branch `cuda-bridge`), on the CPU and on NVIDIA GPUs through CUDA kernels of the Wilson and Möbius operators taken from Bridge++. On the device the fine operators, the prolongation (stored in half precision unless asked otherwise), the coarse operators and the coarse Krylov solvers all stay on the GPU. The cycle above the coarse lattice runs in single precision on operators converted on the device from the links of the double-precision ones; the outer solver is in double precision. The coarse solvers are BiCGStab for the quark factor and MINRES on $\mathcal{H}_{\mathrm{c}}$ for the Pauli–Villars factor (Sec. VI), both stopped at $|r|^2/|b|^2 < 10^{-2}$ unless stated, on the even coarse sites. The numbers below come from the parameter-file program `ruttice` and from the examples `preconditioned_fov` (CPU) and `preconditioned_fov_device` (GPU), whose options are given with each result. The eigensolvers (thick-restart Lanczos, Krylov–Schur Arnoldi [28] with left eigenvectors, and $\sigma_{\min}(z - X)$ by Lanczos on $(z - X)^{\dagger}(z - X)$) are in the module `eigen` and are checked against Bridge++'s eigensolver test.

XI. Numerical results

A. Setups

Timings are from an NVIDIA A30; an RTX 3080 was used for iteration counts and spectral quantities only. Table 2 lists the setups. The smoother interval is not chosen from the operator: setup S needs $\kappa_s \approx 10$ and setup B $\kappa_s \approx 30$, and the wrong one makes the solve stagnate (Sec. G).

TABLE 2. The two setups and the multigrid of the mass scan. All runs: Möbius $(b, c) = (1.5, 0.5)$, $L_s = 8$, $m_{\mathrm{PV}} = 1$, $\alpha_{\mathrm{PV}} = 1$, boundary $(1, 1, 1, -1)$, $\alpha = 0.6$ unless stated. Configurations: the v2 quenched test configurations on 12^4, 16^4 and $24^3 \times 48$ (the 12^4 one has one unpaired near-zero mode, $Q = +1$).

	setup S (stout)	setup B (bare)
M_5	1.0	1.8
gauge smearing	stout, $\rho = 0.1$, 3 steps	none
Chebyshev smoother ν, κ_s	8, 10	12, 30
multigrid (both setups)		
aggregates	4^4 ($24^3 \times 48$: levels below with super-blocks 2, 3)	
test vectors	32 (16 filtered low modes of $\tilde{A}_R^\dagger \tilde{A}_R$, lifted to $\{s, \hat{t}\}$)	
setup	filter 2 rounds; bootstrap 2 rounds of 4 FGMRES steps	
coarse solves	single precision, $\lvert r\rvert^2/\lvert b\rvert^2 < 10^{-2}$	
outer solver	GCR(2) on $\tilde{A}_L$ to $\lvert r\rvert^2/\lvert b\rvert^2 < 10^{-18}$	

B. Comparison with CGNE

Table 3 is the mass scan on $24^3 \times 48$. The CGNE iterations grow by a factor of 21 (S) and 18 (B) from $m = 0.1$ to $m = 0.001$, the multigrid iterations by 2 (S) and 3.7 (B). At $m = 0.001$ the multigrid solve is 4.7 times (S) and 6.0 times (B) faster; with the setup of about 390 s it pays for itself after 4 (S) and 3.4 (B) right-hand sides, i.e. within one point-source propagator (12 columns) for $m \leq 0.005$ (S) and $m \leq 0.01$ (B). For $m \geq 0.05$ in setup S the multigrid is not faster at all. The multigrid needs three times the device memory of CGNE; on the 24 GB A30 this excludes GCR(k) with $k \geq 8$ (out of memory after the setup). Neff's $\alpha = 0.6$ against $\alpha = 1$ in the left form: equal for S, one to three iterations fewer for B. Runs of the same parameter file with two builds (`535c9ef`, `ca1e892`) differ by up to two iterations in three of eight cases; the cause (code change or non-determinism of the device reductions) is not established.

Mixed-precision CGNE. For a comparison on equal terms the reference solver also exists in mixed precision (defect correction in double precision, CG in single precision on the normal equation of $C_e^{-1}S$). A first scan, setup S on $24^3 \times 48$ (double: iterations / s; mixed: iterations / s): $m = 0.1$: 173/5.7, 200/4.7; 0.05: 291/9.5, 341/7.8; 0.02: 644/21.0, 1197/26.9; 0.01: 1166/37.9, 2079/46.4; 0.005: 2067/67.1, 3432/76.4. These runs overshot the tolerance (final $\lvert D\psi - \eta\rvert^2$ of 10^{-19} to 10^{-26}), which has since been corrected; on 12^4 setup S, $m = 0.01$, the corrected scheme needs 301 single-precision iterations against 212 in double. [TODO: the corrected mixed-precision scan, both setups, all masses.]

C. Field of values

The Hermitian part of $\tilde{A}_R\mathcal{M}$ was computed with $\mathcal{M}$ made a fixed linear map: both coarse factors solved to $\lvert r\rvert^2/\lvert b\rvert^2 < 10^{-20}$ in double precision with the exact adjoints, and $\mathcal{M}^\dagger$ built with the same $\mathcal{P}$ and the same smoother interval, so that $\langle z, \tilde{A}_R\mathcal{M}w\rangle$ and $\langle \mathcal{M}^\dagger\tilde{A}_R^\dagger z, w\rangle$ agree to 10^{-10}. The cycle was also taken apart: the smoother alone, $S = p(\tilde{A}_R)$, the coarse correction alone, E of Eq. (5.2), and the cycle $\mathcal{M}$ (Table 4).

On $4^3 \times 8$ the field of values of $\tilde{A}_R\mathcal{M}$ lies in the right half-plane; on 12^4 and 16^4 it does not. The smoother is harmless ($\lVert I - \tilde{A}_R S\rVert \approx 1$); the factor that amplifies is the coarse correction, $\lVert I - \tilde{A}_R E\rVert$ of order 10^3 to 10^4 against 18.5 on $4^3 \times 8$: the factorised correction, exact on the vectors it is built from, is a strongly oblique projection. The bound of Ref. [21] for GCR therefore does not apply (open problem 5). Replacing E by the Galerkin correction of $\tilde{A}_R$ itself, $\mathcal{P}(\mathcal{P}^\dagger\tilde{A}_R\mathcal{P})^{-1}\mathcal{P}^\dagger$ (built densely), brings the field of values of setup B into the right half-plane, $[0.121, 1.99]$, and reduces $\lVert I - \tilde{A}_R\mathcal{M}\rVert$ to 1.84 (setup S: $[-3.62, 13.1]$, 16.8); but GCR(8) to 10^{-16} then needs more iterations, 16 against 13 (B) and 12 against

TABLE 3. Mass scan on $24^3 \times 48$, A30, one random right-hand side to $|r|^2/|b|^2 < 10^{-18}$ on the even-odd system: even-odd CGNE in double precision against the multigrid (GCR(2) on $\tilde{A}_L$), $\alpha = 0.6$ for both and $\alpha = 1$ for the multigrid; iterations, solve time in seconds, coarse-operator applications, and N_{rhs}, the number of right-hand sides at which the multigrid setup (376–394 s) is paid back, $t_{\mathrm{setup}}/(t_{\mathrm{CGNE}} - t_{\mathrm{MG}})$. Final $|D\psi - \eta|^2/|\eta|^2$: 10^{-17}–3×10^{-16} (CGNE), 2×10^{-19}–3×10^{-18} (MG). Device memory peak: CGNE 6.9 GB, MG 19.9 GB ($\alpha = 0.6$). RutticeQCD `535c9ef`.

	CGNE		MG, $\alpha = 0.6$			MG, $\alpha = 1$		
m	it	s	it	s	coarse	it	s	N_{rhs}
setup S ($M_5 = 1.0$, stout)								
0.1	173	5.7	13	11.2	1160	15	12.1	—
0.05	291	9.5	13	9.5	734	15	11.0	—
0.02	644	21.0	14	10.4	810	14	10.5	36
0.01	1166	37.9	15	11.7	1006	15	11.9	15
0.005	2067	67.1	18	13.7	1134	18	14.8	7.2
0.002	3419	110.9	23	19.6	1942	22	20.3	4.3
0.001	3602	116.8	26	24.9	2875	27	27.1	4.3
setup B ($M_5 = 1.8$)								
0.1	231	7.6	7	5.5	190	9	6.9	184
0.05	402	13.1	10	8.1	315	12	9.5	78
0.02	857	27.8	13	10.5	428	15	12.2	22
0.01	1447	46.9	15	12.5	568	17	14.1	11
0.005	2299	74.5	18	15.1	722	21	17.9	6.6
0.002	3473	112.5	23	19.8	1056	25	21.5	4.2
0.001	4228	137.0	26	22.8	1278	27	23.4	3.4

10 (S). The field of values is not what decides the iteration count.

D. Spectrum

Although the field of values reaches far into the left half-plane, every eigenvalue lies at $\mathrm{Re}\,\lambda \geq 0.78$ (Table 5), with moderate condition numbers. Setup S has one (12^4) or two (16^4) isolated outliers of large modulus that grow with the volume; all other eigenvalues lie within 0.43–0.47 of 1. The left and right forms have the same spectrum to the precision shown. At $m = 0.005$ on 16^4 the lowest real parts are 0.802 (S) and 0.790 (B).

E. Pseudospectra and a bound for GMRES

For a non-normal operator the eigenvalues alone do not bound the convergence of GMRES; the ε-pseudospectrum, the set where $\sigma_{\min}(z - X) < \varepsilon$, does [13, 29]. If a contour Γ of length L encloses the spectrum of X but not the origin and $\sigma_{\min}(z - X) \geq \varepsilon$ on it, then for any polynomial p_k with $p_k(0) = 1$,

$$\frac{\|r_k\|}{\|r_0\|} \leq \|p_k(X)\| \leq \frac{L}{2\pi\varepsilon} \max_{z \in \Gamma} |p_k(z)|. \tag{11.1}$$

An isolated outlier λ_o outside Γ is removed by the factor $(1 - z/\lambda_o)$ of p_k: then $p_k(X) = p_k(X)(I - P_o)$ with P_o the spectral projector, $\|(z - X)^{-1}(I - P_o)\| \leq \kappa_o/\varepsilon$, and with a circle $|z - c| = r$ and $p_k(z) = (1 - z/\lambda_o)(1 - z/c)^{k-1}$,

$$\frac{\|r_k\|}{\|r_0\|} \leq \kappa_o \, \frac{r}{\varepsilon} \, \frac{|\lambda_o - c| + r}{|\lambda_o|} \Big(\frac{r}{c}\Big)^{k-1}. \tag{11.2}$$

TABLE 4. Real part of the field of values $W(X)$ and $\|I - X\|_2$ for $X = \tilde{A}_R S$, $\tilde{A}_R E$, $\tilde{A}_R \mathcal{M}$ (exact coarse solves). $m = 0.05$ on $4^3 \times 8$ (2^4 aggregates, 8 test vectors), $m = 0.001$ otherwise; 16 test vectors. `preconditioned_fov_device -parts`.

	$\mathrm{Re}\, W(\tilde{A}_R)$	$\|I - \tilde{A}_R S\|$	$\|I - \tilde{A}_R E\|$	$\mathrm{Re}\, W(\tilde{A}_R \mathcal{M})$	$\|I - \tilde{A}_R \mathcal{M}\|$
$4^3 \times 8$, $M_5 = 1$	$[-5.0 \times 10^{-3}, 2.26]$		18.5	$[0.772, 1.222]$	
12^4 S	$[-4.59 \times 10^{-3}, 1.98]$	1.018	1.94×10^4	$[-23.9, 30.8]$	54.8
12^4 B	$[-5.02 \times 10^{-2}, 4.46]$	1.112	2.86×10^3	$[-1.49, 4.17]$	5.64
16^4 S			1.54×10^4	$[-15.6, 27.1]$	
16^4 B			2.5×10^3	$[-1.42, 4.27]$	

TABLE 5. Eigenvalues of $\tilde{A}\mathcal{M}$ (exact coarse solves) by Krylov–Schur: the smallest real part, the condition numbers $\kappa = 1/|y^\dagger x|$ of the six eigenvalues of smallest real part (left eigenvectors from $\mathcal{M}^\dagger \tilde{A}^\dagger$), the eigenvalues farthest from 1 with their κ, and the radius about 1 of the disc that holds all the others. $m = 0.001$. `preconditioned_fov_device -arnoldi -far 1,0` (`-left` for $\tilde{A}_L$).

	$\min \mathrm{Re}\, \lambda$	κ (6 lowest)	outliers (κ)	others: $\|\lambda - 1| \le$
$4^3 \times 8$, $m = 0.05$, R	0.857	2–4		
$4^3 \times 8$, $m = 0.05$, L	0.854	2.8		0.156 (Ritz)
12^4 S, R	0.868	6.3–14.5	$5.86 + 0.86i$ (11.2)	0.434
12^4 S, L	0.861	7.2–26.2	$5.40 + 0.89i$ (13.9)	0.455
12^4 B, R	0.831	9.4–20.3		0.753
12^4 B, L	0.830			0.747
16^4 S, R	0.794	6.2–12.8	$10.63 + 0.19i$ (4.38), $3.13 + 0.02i$ (4.74)	0.474
16^4 B, R	0.783	13.5–42.9		0.840

On 16^4 setup S has two outliers; the factor κ_o is then replaced by $\|I - P_1 - P_2\| \le 1 + \kappa_1 + \kappa_2$ and there are two factors of the form $(|\lambda_o - c| + r)/|\lambda_o|$. On the real axis between the origin and the spectrum, $\sigma_{\min}(z - \tilde{A}_R \mathcal{M}) \ge 0.024$ (S) and ≥ 0.031 (B) on 12^4 ($z = -0.5 \ldots 0.7$), and 0.107 (S), 0.271 (B) at the origin.

The bound guarantees convergence at a fixed rate, r/c per iteration, without stagnation; it is pessimistic by a factor of 3–5 in iterations (GCR needs 10–17 with exact coarse solves, Sec. F), and its rate is that measured on $24^3 \times 48$, where the median factor per GCR step of $|r|^2$ is 0.23–0.34 (Table 8), i.e. 0.48–0.58 in $|r|$, against $r/c = 0.5$–0.6 for setup S. Setup B is not worse between 12^4 and 16^4 ($\min \sigma$ grows from 0.098 to 0.13); in setup S the outliers grow but their condition numbers fall and ε stays the same.

What is not yet rigorous. (i) $\sigma_{\min}$ is known at 16 points per circle, spaced by more than ε; since $\sigma_{\min}(z - X)$ is 1-Lipschitz in z, a lower bound on the whole circle needs a spacing below ε. [TODO: the dense contours on 12^4 (160 and 96 points), running.] (ii) Lanczos approaches $\sigma^2_{\min}$ from above; a rigorous lower bound needs an inclusion theorem (Kato–Temple) with a bound on the second singular value. (iii) That the disc holds all eigenvalues but the outliers rests on the four eigenvalues farthest from c found by Arnoldi. (iv) The bound is for full GMRES with a fixed linear $\mathcal{M}$ (Sec. F). (v) The FOV and spectra are not available on $24^3 \times 48$: the eigensolvers do not fit in the memory of the A30.

F. Truncated GCR

With the usual coarse solves to 10^{-2}, GCR(2) needs one to five iterations more than the full method (Table 7). With accurate coarse solves, where $\mathcal{M}$ is close to the fixed map of Secs. C–E, GCR(2) stagnates

TABLE 6. $\min_\Gamma \sigma_{\min}(z - \tilde{A}_R\mathcal{M})$ on circles $|z - c| = r$ (16 points each), and the bound (11.1) or (11.2): the factor per iteration r/c, the constant, and the iterations it guarantees for $\|r_k\|/\|r_0\| < 10^{-9}$. Lanczos to 10^{-3} relative in σ^2. `preconditioned_fov_device -pseudo`.

	c	r	$\min \sigma_{\min}$	r/c	constant	k for 10^{-9}
12^4 S	1	0.6	0.0285	0.6	220	53
12^4 S	1	0.5	0.0107	0.5	480	40
12^4 B	1.4	0.93	0.0985	0.66	9.4	57
12^4 B	1.3	0.87	0.0860	0.67	10.1	58
16^4 S	1	0.6	0.0262	0.6	190	53
16^4 S	1	0.55	0.0135	0.55	350	46
16^4 B	1.4	1.05	0.1295	0.75	8.1	80

TABLE 7. GCR(k) iterations to $|r|^2/|b|^2 < 10^{-18}$, 12^4, $m = 0.001$, the multigrid of Table 2 (no levels below the coarse lattice), one random right-hand side; the point source: range over its 12 columns.

	coarse solves	$k = 2$	8	16	32	64
S, $\tilde{A}_L$	10^{-2}	14			12	
S, $\tilde{A}_R$	10^{-2}	12			11	
B, $\tilde{A}_L$	10^{-2}	14			13	
B, $\tilde{A}_R$	10^{-2}	15			14	
S, $\tilde{A}_L$, point source	10^{-2}	15–19			13–14	
S, $\tilde{A}_L$	10^{-8}	no conv.	26	17	17	17
S, $\tilde{A}_L$	10^{-4}	18			16	

($|r|^2 \approx 3 \times 10^{-8}$ after 1000 iterations) while GCR(8) converges in 26 and GCR(≥ 16) in 17. This is what the analysis predicts: the pseudospectra guarantee the full method; a truncated or restarted one has no such guarantee when the field of values contains the origin [21, 30]. The truncated solver works in practice because the inexact coarse solves do not let the correction amplify as E does, an empirical fact without a proof. On $24^3 \times 48$ the full method cannot be run on the A30: GCR(8) and GCR(16) run out of device memory after the setup (GCR(2) peaks at 19.9 of 24 GB, and every further direction is another fine-grid vector).

No run shows a plateau (Table 8); the first step reduces $|r|^2$ by about 10^{-5} in every run.

G. Sensitivity to the parameters

The method is robust within a window of each parameter and fails outside it (Table 9). The smoother interval must suit the setup; the bootstrap is necessary; more accurate coarse solves are not better, and with GCR(2) accurate ones stagnate (Sec. F). Larger setups (smaller aggregates, more test vectors) buy fewer iterations at a setup cost that only many right-hand sides repay. Every stagnation is of GCR(2). [TODO: timings of this table on the A30; the local coherence of the test vectors.]

XII. Open problems

The following are questions of theory and algorithm, not findings.

1. **Exactness for the left form.** Proposition 1 holds for the right form. For $\tilde{A}_L$ (with or without α) the Pauli–Villars lift that would be needed differs from $\hat{t}$ by site-local factors (second remark of Sec. C).

TABLE 8. Residual history of GCR(2), $24^3 \times 48$, $\tilde{A}_L$, the multigrid of Table 2: steps, median and largest factor of $|r|^2$ per step, and the longest stretch of steps with a factor above $1/2$.

setup	m, α	steps	median	worst	slow stretch
S	0.001, 0.6	28	0.315	0.564	2 steps
S	0.001, 1.0	27	0.338	0.479	none
S	0.002, 0.6	23	0.233	0.426	none
S	0.002, 1.0	23	0.253	0.354	none
B	0.001, 0.6	26	0.299	0.357	none
B	0.001, 1.0	28	0.317	0.403	none
B	0.002, 0.6	23	0.246	0.300	none
B	0.002, 1.0	25	0.280	0.347	none

TABLE 9. GCR(2) iterations when one parameter of Table 2 is changed, 12^4, $m = 0.001$, $\tilde{A}_L$, one random right-hand side (base: 14 in both setups); "stag.": not converged in 1000 iterations. Setup times in seconds where they change much.

change	setup S	setup B
aggregates 2^4 / 3^4 / 6^4	9 (89 s) / 12 / 23	9 (61 s) / 11 / 20
test vectors 8 / 16 / 24 / 48	stag. / 27 / 22 / 10 (40 s)	39 / 19 / 16 / 11 (30 s)
smoother $\nu = 4$ / 6 / 8 / 12 / 16	31 / 18 / 14 / 12 / 9	stag. / 84 / 26 / 14 / 13
interval $\kappa_s = 5$ / 10 / 30 / 100	22 / 14 / stag. / stag.	stag. / 31 / 14 / stag.
coarse solves 10^{-1} / 10^{-2} / 10^{-4} / 10^{-8}	23 / 14 / 18 / stag.	19 / 14 / 16 / 17
bootstrap rounds 0 / 1 / 2 / 3	stag. / 16 / 14 / 23	249 / 47 / 14 / 12
right form $\tilde{A}_R$	12	15

Whether a basis can be made exact for both forms at once, and what is lost by using the right-form basis for the left form, is open.

2. **Loss of Γ_5-Hermiticity for $c \neq 0$.** By Eq. (3.4), $\Gamma_5 D \Gamma_5 \neq D^\dagger$ at $O(c)$. MINRES on $\mathcal{H}_c = \mathcal{D}_c^{\mathrm{PV}}\Sigma$, the absolute-value block preconditioner (built from the Hermitian part only), the additive level preconditioner, and the Σ-adjoint in place of the exact adjoint are all exact only for $c = 0$. The exactly Hermitian operator $\Gamma_5 D_-^{-1} D$ [8] involves a non-local inverse. How large the non-Hermitian part is on the coarse space, and how to treat it, is open.

3. **Smoother interval for a non-normal operator.** The interval $[\lambda_{\mathrm{hi}}/\kappa_s, \lambda_{\mathrm{hi}}]$ is set from an estimate of the spectral radius, whereas the setup works with singular values (eigenvalues of $\tilde{A}_R^\dagger \tilde{A}_R$). For a non-normal $\tilde{A}$ these need not agree. Which part of the low end is left to the smoother after the coarse correction, and whether an ellipse with complex foci should replace the real interval, are open.

4. **Even-odd smoother and coarse space.** Boyle and Yamaguchi [18] found that a red-black smoother combined with the HDCR coarse correction was not effective, and attributed it to the eigenvectors of the red-black operator and of $\gamma_5 R_5 D$ not being aligned. In the present method the smoother and the correction both approximate $\tilde{A}^{-1}$ and the coarse space is built from low modes of $\tilde{A}_R$, so that particular mismatch is absent by construction; but the coarse operators are Galerkin projections of D and D_{PV}, not of $\tilde{A}$. Whether the even-odd outer system is preferable to the full-lattice A' for this method has not been established.

5. **Field of values versus real parts of eigenvalues.** Section II shows $\mathrm{Re}\,\lambda > 0$ for A'. The GCR bound [21] needs a positive definite Hermitian part, i.e. a field of values in the right half-plane, which is a stronger condition. Whether it holds for A', $\tilde{A}_R$ and $\tilde{A}_L$, and whether a shift of the kind used in Ref. [20] is needed, is not known. For the preconditioned operator $\tilde{A}_R\mathcal{M}$ it does not hold on 12^4 and 16^4 (Sec. C); the spectrum stays at $\mathrm{Re}\,\lambda \geq 0.78$ and the pseudospectra give a bound for full GMRES (Sec. E), but not for truncated GCR (Sec. F).

6. **Outer solver with a varying preconditioner.** Because the coarse solves are truncated, $\mathcal{M}$ is nonlinear and depends on the right-hand side. FGMRES and flexible GCR allow this; BiCGStab does not, and its behaviour with $\mathcal{M}$ as preconditioner has no theoretical justification. The truncated GCR(2) used in practice converges with coarse solves to 10^{-2} and stagnates with accurate ones (Sec. F); why the inexact correction behaves better is open.

7. **Solving with $D^\dagger$.** The method is built for $D\psi = \eta$. A solve with $D^\dagger$, as needed for the normal equations in HMC, requires the dual correction (the two factors and the roles of the lifts exchanged), and its basis and exactness have not been derived.

References

[1] R. Narayanan and H. Neuberger, *Infinitely many regulator fields for chiral fermions*, Phys. Lett. B **302** (1993) 62 [hep-lat/9212019].

[2] H. Neuberger, *Exactly massless quarks on the lattice*, Phys. Lett. B **417** (1998) 141 [hep-lat/9707022].

[3] D. B. Kaplan, *A method for simulating chiral fermions on the lattice*, Phys. Lett. B **288** (1992) 342 [hep-lat/9206013].

[4] Y. Shamir, *Chiral fermions from lattice boundaries*, Nucl. Phys. B **406** (1993) 90 [hep-lat/9303005].

[5] V. Furman and Y. Shamir, *Axial symmetries in lattice QCD with Kaplan fermions*, Nucl. Phys. B **439** (1995) 54 [hep-lat/9405004].

[6] R. C. Brower, H. Neff and K. Orginos, *Möbius fermions: improved domain wall chiral fermions*, Nucl. Phys. B Proc. Suppl. **140** (2005) 686 [hep-lat/0409118].

[7] R. C. Brower, H. Neff and K. Orginos, *Möbius fermions*, Nucl. Phys. B Proc. Suppl. **153** (2006) 191 [hep-lat/0511031].

[8] R. C. Brower, H. Neff and K. Orginos, *The Möbius domain wall fermion algorithm*, Comput. Phys. Commun. **220** (2017) 1 [arXiv:1206.5214].

[9] M. Lüscher, *Local coherence and deflation of the low quark modes in lattice QCD*, JHEP **07** (2007) 081 [arXiv:0706.2298].

[10] J. Brannick, R. C. Brower, M. A. Clark, J. C. Osborn and C. Rebbi, *Adaptive multigrid algorithm for lattice QCD*, Phys. Rev. Lett. **100** (2008) 041601 [arXiv:0707.4018].

[11] A. Frommer, K. Kahl, S. Krieg, B. Leder and M. Rottmann, *Adaptive aggregation-based domain decomposition multigrid for the lattice Wilson–Dirac operator*, SIAM J. Sci. Comput. **36** (2014) A1581 [arXiv:1303.1377].

[12] I. Kanamori, K.-I. Ishikawa and H. Matsufuru, *Object-oriented implementation of algebraic multigrid solver for lattice QCD on SIMD architectures and GPU clusters*, in *Computational Science and Its Applications – ICCSA 2021*, Lecture Notes in Computer Science **12953**, Springer (2021), pp. 218–233, doi:10.1007/978-3-030-86976-2_15 [arXiv:2111.05012].

[13] N. M. Nachtigal, S. C. Reddy and L. N. Trefethen, *How fast are nonsymmetric matrix iterations?*, SIAM J. Matrix Anal. Appl. **13** (1992) 778.

[14] P. A. Boyle, *Multiple right hand side multigrid for domain wall fermions with a multigrid preconditioned block conjugate gradient algorithm*, arXiv:2409.03904.

[15] S. D. Cohen, R. C. Brower, M. A. Clark and J. C. Osborn, *Multigrid algorithms for domain-wall fermions*, PoS **LATTICE2011** (2011) 030 [arXiv:1205.2933].

[16] P. A. Boyle, *Hierarchically deflated conjugate gradient*, arXiv:1402.2585.

[17] A. Yamaguchi and P. A. Boyle, *Hierarchically deflated conjugate residual*, PoS **LATTICE2016** (2016) 374.

[18] P. A. Boyle and A. Yamaguchi, *Comparison of domain wall fermion multigrid methods*, arXiv:2103.05034.

[19] R. C. Brower, M. A. Clark, D. Howarth and E. S. Weinberg, *Multigrid for chiral lattice fermions: domain wall*, Phys. Rev. D **102** (2020) 094517 [arXiv:2004.07732].

[20] P. A. Boyle and P. Oare (with S. Yamamoto, C. Jung and C. Lehner), *Algorithms for domain wall fermions*, talk at Lattice 2026.

[21] S. C. Eisenstat, H. C. Elman and M. H. Schultz, *Variational iterative methods for nonsymmetric systems of linear equations*, SIAM J. Numer. Anal. **20** (1983) 345.

[22] I. Kanamori, W.-L. Chen and H. Matsufuru, *Spectrum of preconditioned Moebius domain-wall operators*, PoS **LATTICE2024** (2025) 414.

[23] H. Neff, *A better conditioned domain wall operator*, arXiv:1501.04950.

[24] W.-L. Chen, I. Kanamori, H. Matsufuru and H. Neff, *Accelerating iterative linear equation solver using modified domain-wall fermion matrix in lattice QCD simulations*, in *Computational Science and Its Applications – ICCSA 2025 Workshops*, Lecture Notes in Computer Science, Springer (2025), pp. 54–69, doi:10.1007/978-3-031-97606-3_5.

[25] E. Vecharynski and A. V. Knyazev, *Absolute value preconditioning for symmetric indefinite linear systems*, SIAM J. Sci. Comput. **35** (2013) A696.

[26] T. A. Manteuffel, *The Tchebychev iteration for nonsymmetric linear systems*, Numer. Math. **28** (1977) 307.

[27] Y. Saad, *Iterative Methods for Sparse Linear Systems*, 2nd ed., SIAM, Philadelphia (2003).

[28] G. W. Stewart, *A Krylov–Schur algorithm for large eigenproblems*, SIAM J. Matrix Anal. Appl. **23** (2001) 601.

[29] L. N. Trefethen and M. Embree, *Spectra and Pseudospectra: The Behavior of Nonnormal Matrices and Operators*, Princeton University Press (2005).

[30] M. Embree, *The tortoise and the hare restart GMRES*, SIAM Rev. **45** (2003) 259.

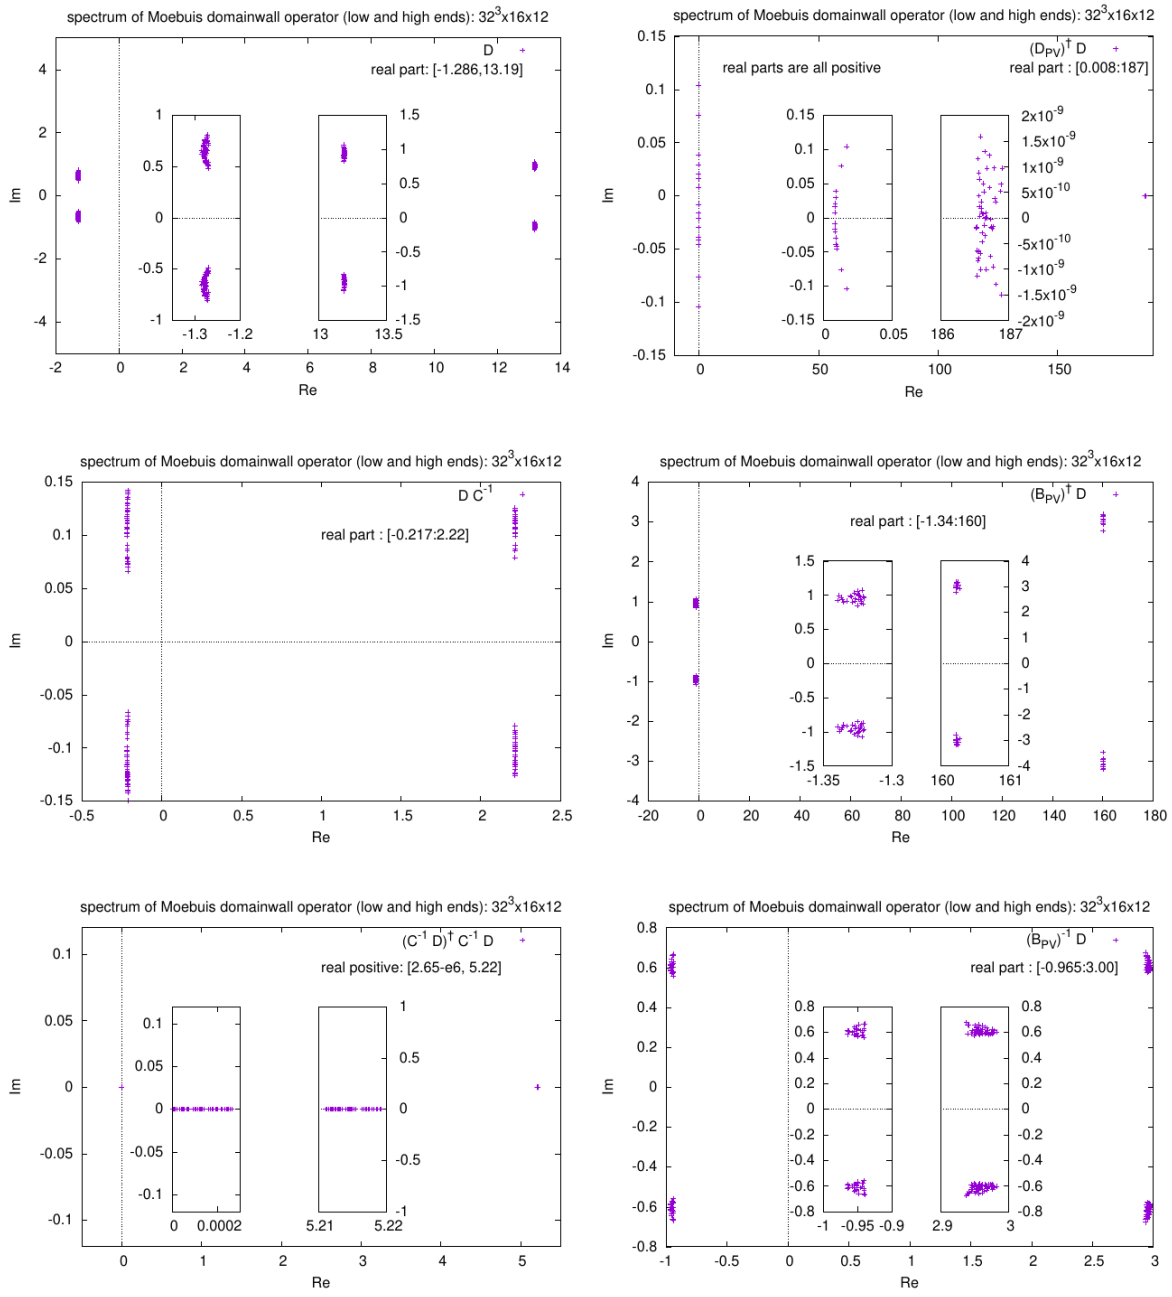


FIG. 2. Low and high ends of the spectrum of the Möbius operator on a JLQCD $2+1$-flavour configuration at the physical point, with preconditioners. Left column: D, DC^{-1}, $(C^{-1}D)^\dagger C^{-1}D$. Right column: $D_{\rm PV}^\dagger D$, $B_{\rm PV}^\dagger D$, $B_{\rm PV}^{-1}D$, where $B_{\rm PV}$ is the block-diagonal approximation of $D_{\rm PV}$ (inverted approximately in the last panel). The text of Ref. [22] quotes $48^3 \times 16$ with $L_s = 12$; the panel titles read $32^3 \times 16 \times 12$. Reproduced from Ref. [22], Fig. 3.

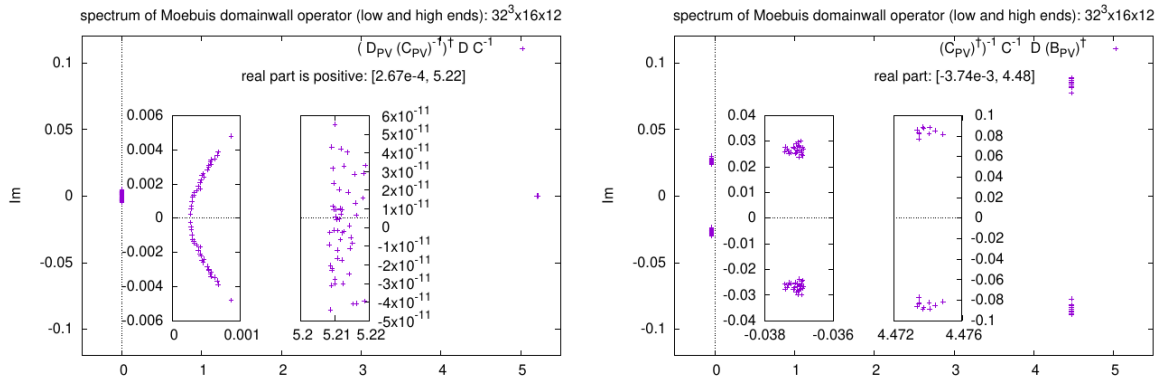


FIG. 3. Left: spectrum of $A' = (D_{\rm PV}C_{\rm PV}^{-1})^\dagger DC^{-1}$, Eq. (3.7), on the configuration of Fig. 2; all real parts are positive. Right: the same with $D_{\rm PV}$ replaced by its block-diagonal approximation $B_{\rm PV}$; negative real parts reappear. Lattice size as in Fig. 2 (the text of the source quotes $48^3 \times 16$, $L_s = 12$; the panel titles read $32^3 \times 16 \times 12$). Reproduced from Ref. [22], Fig. 4.