The factorized Pauli–Villars multigrid of RutticeQCD, from the solver to HMC: a step-by-step walkthrough

Wei-Lun Chen

October 23, 2026 (internal note)

This is an internal walkthrough of the multigrid method for the Möbius domain-wall operator as it is implemented in RutticeQCD, written so that every formula can be followed line by line. It starts from the operator, goes through the solver of $D\psi = \eta$, the setup of the multigrid and the solves with $D^\dagger$, and ends with the solves of the HMC force. Nothing here is new with respect to the MG note (`docs/mg_background/MDWF_MULTIGRID_NOTE.tex`); the note states the results, this walkthrough derives them step by step. The last section maps each step to the code.

Contents

1 The Möbius domain-wall operator **2**
- 1.1 Ingredients . . . 2
- 1.2 The operator in Brower–Neff–Orginos form . . . 2
- 1.3 Splitting off the fifth-dimension hop . . . 3
- 1.4 The site-diagonal part C and its inverse . . . 3
- 1.5 Neff's α . . . 4

2 The even-odd system **4**
- 2.1 Blocks . . . 4
- 2.2 The Schur complement . . . 4
- 2.3 $S^{-1} = \pi D^{-1} \iota$ . . . 5
- 2.4 The Pauli–Villars preconditioned forms . . . 5

3 The outer solver: flexible GCR **6**

4 The two-grid cycle **6**
- 4.1 The Chebyshev smoother, derived . . . 6
- 4.2 The factorized coarse correction . . . 7
- 4.3 Aggregation and the chirality split . . . 8
- 4.4 Orthonormalization per block: CholeskyQR2 . . . 8
- 4.5 The basis $\{s, \hat{t}\}$ and why the correction is exact . . . 9
- 4.6 The coarse operators by probing . . . 10
- 4.7 The coarse solves . . . 10

5 Setup **11**
- 5.1 Chebyshev-filtered subspace iteration . . . 11

5.2 Bootstrap . . . 11
5.3 Cost and memory . . . 11

6 Solving with $D^\dagger$ **11**
6.1 $\Gamma_5 D \Gamma_5 \neq D^\dagger$ . . . 11
6.2 The exact relation . . . 12
6.3 $\Gamma_5 D^{-1} \Gamma_5$ as a preconditioner of $D^\dagger$ . . . 12
6.4 On the even sites, step by step . . . 13

7 HMC **13**
7.1 The even-odd Ratio action . . . 13
7.2 $(K^\dagger K)^{-1}$ with the multigrid: two solves . . . 14
7.3 Neff's α in HMC . . . 15
7.4 What is kept along a trajectory . . . 15
7.5 Memory (24^4, $m = 0.001$, 20 vectors, levels below, A30) . . . 15
7.6 What the measurements say so far . . . 15

8 Map to the code **16**

1 The Möbius domain-wall operator

1.1 Ingredients

The field ψ lives on the four-dimensional lattice and carries a fifth index $s = 0, \dots, L_s - 1$; each ψ_s is a Dirac field. The chiral projectors are

$$P_\pm = \frac{1 \pm \gamma_5}{2}, \qquad P_\pm^2 = P_\pm, \qquad P_+ P_- = 0, \qquad P_+ + P_- = 1. \tag{1.1}$$

The Wilson hopping term and the Wilson operator of mass $-M_5$ are

$$(K\psi)(x) = \sum_{\mu=1}^{4} \Big[(1 - \gamma_\mu) U_\mu(x) \psi(x + \hat{\mu}) + (1 + \gamma_\mu) U_\mu^\dagger(x - \hat{\mu}) \psi(x - \hat{\mu}) \Big], \tag{1.2}$$

$$W = (4 - M_5) - \tfrac{1}{2} K. \tag{1.3}$$

W acts on every slice s alike and does not mix slices.

1.2 The operator in Brower–Neff–Orginos form

With b, c the Möbius parameters, define

$$D_+ = 1 + bW, \qquad D_- = cW - 1. \tag{1.4}$$

The operator at quark mass m acts slice by slice as

$$(D\psi)_s = D_+ \psi_s + D_- \big(P_- \psi_{s+1} + P_+ \psi_{s-1} \big), \qquad \psi_{L_s} \equiv -m\psi_0, \qquad \psi_{-1} \equiv -m\psi_{L_s - 1}. \tag{1.5}$$

For $L_s = 4$ the matrix in s is

$$D = \begin{pmatrix} D_+ & D_- P_- & 0 & -m D_- P_+ \\ D_- P_+ & D_+ & D_- P_- & 0 \\ 0 & D_- P_+ & D_+ & D_- P_- \\ -m D_- P_- & 0 & D_- P_+ & D_+ \end{pmatrix}. \tag{1.6}$$

The physical quark field is $q = P_- \psi_0 + P_+ \psi_{L_s - 1}$; the runs use $b = 1.5$, $c = 0.5$, $L_s = 8$.

1.3 Splitting off the fifth-dimension hop

Define the hop in the fifth dimension, with a minus sign,

$$(Y\psi)_s = -P_-\psi_{s+1} - P_+\psi_{s-1} \qquad \text{(same wrap-around)}. \tag{1.7}$$

Comparing with Eq. (1.5), the off-diagonal part of D is $-D_-Y$, with D_- on the *left*:

$$D = D_+ - D_-Y. \tag{1.8}$$

Now insert $D_\pm$ and then W, one substitution per line:

$$\begin{aligned}
D &= (1 + bW) - (cW - 1)Y && (1.9)\\
&= 1 + bW - cWY + Y && (1.10)\\
&= 1 + Y + W\,(b - cY) && (1.11)\\
&= 1 + Y + \Big[(4 - M_5) - \tfrac{1}{2}K\Big](b - cY) && (1.12)\\
&= \underbrace{\big[1 + b(4 - M_5)\big]}_{d_+} + \underbrace{\big[1 - c(4 - M_5)\big]}_{d_-} Y - \tfrac{1}{2}K\,(b - cY) && (1.13)\\
&= C - \tfrac{1}{2}K\,(b - cY), \qquad C \equiv d_+ + d_-Y. && (1.14)
\end{aligned}$$

C contains no four-dimensional hop: it acts within one site, as an $L_s \times L_s$ matrix in s. All the gauge field sits in $-\frac{1}{2}K(b - cY)$. The Pauli–Villars operator is the same operator at unit mass,

$$D_{\mathrm{PV}} = D\big|_{m=1}, \qquad C_{\mathrm{PV}} = C\big|_{m=1}. \tag{1.15}$$

1.4 The site-diagonal part C and its inverse

Y maps $P_+\psi$ to $P_+\psi$ and $P_-\psi$ to $P_-\psi$, so C splits into a block on $\psi^+ = P_+\psi$ and one on $\psi^- = P_-\psi$. On ψ^+, Y brings ψ^+_{s-1} to slice s with the factor -1, and across the wall the factor is $-(-m) = +m$:

$$(C\psi)^+_s = \begin{cases} d_+\psi^+_0 + m\,d_-\psi^+_{L_s-1}, & s = 0,\\ d_+\psi^+_s - d_-\psi^+_{s-1}, & s \geq 1. \end{cases} \tag{1.16}$$

To solve $C^+x = r$, take the rows $s \geq 1$ first:

$$d_+x_s - d_-x_{s-1} = r_s \quad \Longrightarrow \quad x_s = \frac{r_s + d_-x_{s-1}}{d_+}. \tag{1.17}$$

Each x_s is then affine in the unknown x_0. With $q = d_-/d_+$ we claim

$$x_s = p_s + q^s x_0, \qquad p_0 = 0, \qquad p_s = \frac{r_s + d_-p_{s-1}}{d_+}. \tag{1.18}$$

Check by induction:

$$\begin{aligned}
x_s &= \frac{r_s + d_-x_{s-1}}{d_+} && (1.19)\\
&= \frac{r_s + d_-(p_{s-1} + q^{s-1}x_0)}{d_+} && (1.20)\\
&= \frac{r_s + d_-p_{s-1}}{d_+} + \frac{d_-}{d_+}q^{s-1}x_0 && (1.21)\\
&= p_s + q^s x_0. && (1.22)
\end{aligned}$$

The remaining row $s = 0$ fixes x_0:

$$d_+ x_0 + m\, d_- x_{L_s-1} = r_0 \tag{1.23}$$
$$d_+ x_0 + m\, d_- (p_{L_s-1} + q^{L_s-1} x_0) = r_0 \tag{1.24}$$
$$(d_+ + m\, d_- q^{L_s-1}) x_0 = r_0 - m\, d_- p_{L_s-1} \tag{1.25}$$
$$x_0 = \frac{r_0 - m\, d_- p_{L_s-1}}{d_+ + m\, d_- q^{L_s-1}}. \tag{1.26}$$

The block on ψ^- is the same with $s \to L_s - 1 - s$. So C^{-1} costs $O(L_s)$ per site, is exact, and does not depend on the gauge field. $C^\dagger$ has the transposed blocks and is inverted the same way from the other end.

1.5 Neff's α

Neff's operator is $D^{(\alpha)} = DQ$ with the site-local

$$Q_0 = P_- + \alpha P_+, \qquad Q_s = \alpha \ (0 < s < L_s - 1), \qquad Q_{L_s-1} = P_+ + \alpha P_-. \tag{1.27}$$

Q is Hermitian and Q^{-1} is the same with $\alpha \to 1/\alpha$. If $D^{(\alpha)}\psi^{(\alpha)} = \eta$, then

$$\eta = D^{(\alpha)}\psi^{(\alpha)} = DQ\psi^{(\alpha)} \quad \Longrightarrow \quad \psi = Q\psi^{(\alpha)} \text{ solves } D\psi = \eta. \tag{1.28}$$

Q leaves $P_-\psi_0$ and $P_+\psi_{L_s-1}$ alone, so the quark propagator is unchanged.

2 The even-odd system

2.1 Blocks

Order the sites even first. K connects nearest neighbours, which have opposite parity, so in Eq. (1.14) the diagonal blocks come from C alone:

$$D = \begin{pmatrix} C_e & D_{eo} \\ D_{oe} & C_o \end{pmatrix}, \qquad D_{pq} = -\tfrac{1}{2} K_{pq}(b - cY) \quad (p \neq q). \tag{2.1}$$

C_e and C_o are the same matrix C on the even and the odd sites.

2.2 The Schur complement

Solve $D\psi = \eta$ block by block. The odd rows:

$$D_{oe}\psi_e + C_o\psi_o = \eta_o \tag{2.2}$$
$$\psi_o = C_o^{-1}(\eta_o - D_{oe}\psi_e). \tag{2.3}$$

Insert into the even rows:

$$C_e\psi_e + D_{eo}\psi_o = \eta_e \tag{2.4}$$
$$C_e\psi_e + D_{eo}C_o^{-1}(\eta_o - D_{oe}\psi_e) = \eta_e \tag{2.5}$$
$$(C_e - D_{eo}C_o^{-1}D_{oe})\psi_e = \eta_e - D_{eo}C_o^{-1}\eta_o \tag{2.6}$$
$$S\psi_e = \eta'_e, \qquad S \equiv C_e - D_{eo}C_o^{-1}D_{oe}, \quad \eta'_e \equiv \eta_e - D_{eo}C_o^{-1}\eta_o. \tag{2.7}$$

2.3 $S^{-1} = \pi D^{-1} \iota$

Let ι put an even-site field on the whole lattice (zero on the odd sites) and π take the even part:

$$\iota y = \begin{pmatrix} y \\ 0 \end{pmatrix}, \qquad \pi \begin{pmatrix} \psi_e \\ \psi_o \end{pmatrix} = \psi_e. \tag{2.8}$$

For $\eta = \iota \eta_e$ we have $\eta_o = 0$, so $\eta'_e = \eta_e$ and Eq. (2.7) reads $S\psi_e = \eta_e$. Hence

$$S^{-1}\eta_e = \psi_e = \pi D^{-1} \iota \eta_e \quad \Longrightarrow \quad S^{-1} = \pi D^{-1} \iota. \tag{2.9}$$

The same for $D^\dagger$: its block form is $\begin{pmatrix} C_e^\dagger & D_{oe}^\dagger \\ D_{eo}^\dagger & C_o^\dagger \end{pmatrix}$, and elimination gives $C_e^\dagger - D_{oe}^\dagger (C_o^\dagger)^{-1} D_{eo}^\dagger = S^\dagger$, so $S^{-\dagger} = \pi D^{-\dagger} \iota$.

2.4 The Pauli–Villars preconditioned forms

S is not suitable for a Krylov method: its spectrum surrounds the origin. Following Kanamori, Chen and Matsufuru we multiply from the left by the adjoint of the Pauli–Villars factor and divide the site part out. Two placements of C_e give two forms:

$$\tilde{A}_R = (S_{\mathrm{PV}} C_{\mathrm{PV},e}^{-1})^\dagger \, (S C_e^{-1}), \tag{2.10}$$

$$\tilde{A}_L = (C_{\mathrm{PV},e}^{-1} S_{\mathrm{PV}})^\dagger \, (C_e^{-1} S). \tag{2.11}$$

For $m = 1$ both become $X^\dagger X$, Hermitian positive; for $m < 1$ they are perturbations of it, with all eigenvalues in the right half-plane.

Right form, step by step. Start from $S\psi_e = \eta'_e$ and substitute $\psi_e = C_e^{-1} y$:

$$S C_e^{-1} y = \eta'_e \tag{2.12}$$

$$(S_{\mathrm{PV}} C_{\mathrm{PV},e}^{-1})^\dagger S C_e^{-1} y = (S_{\mathrm{PV}} C_{\mathrm{PV},e}^{-1})^\dagger \eta'_e \tag{2.13}$$

$$\tilde{A}_R \, y = \rho_R, \qquad \rho_R \equiv (S_{\mathrm{PV}} C_{\mathrm{PV},e}^{-1})^\dagger \eta'_e, \tag{2.14}$$

$$\psi_e = C_e^{-1} y, \qquad \psi_o = C_o^{-1}(\eta_o - D_{oe}\psi_e). \tag{2.15}$$

Left form with α in D only. For DQ the blocks are $C_e Q_e$, $D_{eo} Q_o$, $D_{oe} Q_e$, $C_o Q_o$, and the Schur complement is

$$S^{(\alpha)} = C_e Q_e - D_{eo} Q_o (C_o Q_o)^{-1} D_{oe} Q_e \tag{2.16}$$

$$= C_e Q_e - D_{eo} Q_o Q_o^{-1} C_o^{-1} D_{oe} Q_e \tag{2.17}$$

$$= (C_e - D_{eo} C_o^{-1} D_{oe}) Q_e \tag{2.18}$$

$$= S Q_e. \tag{2.19}$$

So the left Schur factor of DQ is

$$K^{(\alpha)} = (C_e Q_e)^{-1} S Q_e = Q_e^{-1} C_e^{-1} S Q_e = Q_e^{-1} K Q_e, \qquad K \equiv C_e^{-1} S, \tag{2.20}$$

and the left form becomes $\tilde{A}_L^{(\alpha)} = (C_{\mathrm{PV},e}^{-1} S_{\mathrm{PV}})^\dagger Q_e^{-1} (C_e^{-1} S) Q_e$. In the right form the two Q_e cancel: $S Q_e (C_e Q_e)^{-1} = S C_e^{-1}$, so $\tilde{A}_R$ does not see α.

3 The outer solver: flexible GCR

The outer iteration solves $\tilde{A}y = \rho$ ($\tilde{A} = \tilde{A}_R$ or $\tilde{A}_L$) with a right preconditioner $\mathcal{M}$, the multigrid cycle. Because $\mathcal{M}$ contains Krylov solves stopped by a residual test, it changes from one application to the next, so the method must be flexible: it keeps the preconditioned directions z_i themselves.

One step of GCR(k), from the residual $r_i = \rho - \tilde{A}y_i$:

$$z_i = \mathcal{M} r_i \qquad \text{(one cycle)} \tag{3.1}$$

$$q_i = \tilde{A} z_i \qquad \text{(one application of } \tilde{A}) \tag{3.2}$$

$$q_i \leftarrow q_i - \sum_j \frac{(q_j, q_i)}{(q_j, q_j)} q_j, \quad z_i \leftarrow z_i - \sum_j \frac{(q_j, q_i)}{(q_j, q_j)} z_j \qquad \text{(the last } k \text{ kept pairs)} \tag{3.3}$$

$$\alpha_i = \frac{(q_i, r_i)}{(q_i, q_i)} \tag{3.4}$$

$$y_{i+1} = y_i + \alpha_i z_i, \qquad r_{i+1} = r_i - \alpha_i q_i. \tag{3.5}$$

The coefficient α_i minimizes $\|r_i - \alpha q_i\|$, and the orthogonalization keeps $\tilde{A}z_i = q_i$ true after the update, so $r_{i+1} = \rho - \tilde{A}y_{i+1}$ exactly. With $k = 2$ (GCR(2)) the code keeps $k + 1 = 3$ pairs (z, q) and the residual, seven fields on the even sites, at the price of no optimality over the whole Krylov space.

4 The two-grid cycle

One cycle $y = \mathcal{M}b$ is

$$y \leftarrow p(\tilde{A})\, b \qquad \text{(pre-smoothing)} \tag{4.1}$$

$$r \leftarrow b - \tilde{A}y \tag{4.2}$$

$$y \leftarrow y + e(r) \qquad \text{(coarse correction)} \tag{4.3}$$

$$r \leftarrow b - \tilde{A}y \tag{4.4}$$

$$y \leftarrow y + p(\tilde{A})\, r \qquad \text{(post-smoothing).} \tag{4.5}$$

4.1 The Chebyshev smoother, derived

The smoother is ν steps of Chebyshev iteration on $\tilde{A}x = b$ from $x = 0$. Choose the interval $[\lambda_{\rm lo}, \lambda_{\rm hi}]$ that it damps ($\lambda_{\rm hi}$ slightly above the spectral radius from a power iteration, $\lambda_{\rm lo} = \lambda_{\rm hi}/\kappa_s$) and set

$$\theta = \tfrac{1}{2}(\lambda_{\rm hi} + \lambda_{\rm lo}), \qquad \delta = \tfrac{1}{2}(\lambda_{\rm hi} - \lambda_{\rm lo}), \qquad \sigma = \theta/\delta. \tag{4.6}$$

The residual after k steps is $r_k = R_k(\tilde{A})r_0$ with the scaled Chebyshev polynomial

$$R_k(\lambda) = \frac{T_k\big((\theta - \lambda)/\delta\big)}{T_k(\sigma)}, \qquad R_k(0) = 1, \tag{4.7}$$

which is at most $1/T_k(\sigma)$ in modulus on the interval. Write $t = (\theta - \lambda)/\delta$ and use $T_{k+1}(t) = 2tT_k(t) - T_{k-1}(t)$:

$$T_{k+1}(\sigma) R_{k+1} = 2t\, T_k(\sigma) R_k - T_{k-1}(\sigma) R_{k-1} \tag{4.8}$$

$$R_{k+1} = \rho_k \big[2t\, R_k - \rho_{k-1} R_{k-1}\big], \qquad \rho_k \equiv \frac{T_k(\sigma)}{T_{k+1}(\sigma)}. \tag{4.9}$$

The same recurrence at $t = \sigma$ gives the ρ's:

$$T_{k+1}(\sigma) = 2\sigma T_k(\sigma) - T_{k-1}(\sigma) \tag{4.10}$$

$$\frac{1}{\rho_k} = 2\sigma - \rho_{k-1}, \qquad \rho_0 = \frac{T_0(\sigma)}{T_1(\sigma)} = \frac{1}{\sigma}. \tag{4.11}$$

Now split $2t = 2\sigma - 2\lambda/\delta$ and use $2\sigma = 1/\rho_k + \rho_{k-1}$:

$$R_{k+1} = \rho_k\Big[(2\sigma - \tfrac{2\lambda}{\delta})R_k - \rho_{k-1}R_{k-1}\Big] \tag{4.12}$$

$$= \rho_k\Big[(\tfrac{1}{\rho_k} + \rho_{k-1})R_k - \tfrac{2\lambda}{\delta}R_k - \rho_{k-1}R_{k-1}\Big] \tag{4.13}$$

$$= R_k + \rho_k\rho_{k-1}(R_k - R_{k-1}) - \frac{2\rho_k}{\delta}\lambda R_k. \tag{4.14}$$

Apply to r_0, with $\lambda \to \tilde{A}$:

$$r_{k+1} - r_k = \rho_k\rho_{k-1}(r_k - r_{k-1}) - \frac{2\rho_k}{\delta}\tilde{A}r_k. \tag{4.15}$$

Since $r = b - \tilde{A}x$, a step $x_{k+1} = x_k + d_k$ changes the residual by $-\tilde{A}d_k$. Matching the two sides,

$$d_k = \rho_k\rho_{k-1}\, d_{k-1} + \frac{2\rho_k}{\delta}\, r_k. \tag{4.16}$$

The first step: $R_1 = T_1(t)/T_1(\sigma) = t/\sigma = 1 - \lambda/\theta$, so $x_1 = b/\theta$ and $d_0 = r_0/\theta$. These are the updates of the code (`BlasChebyshevStep`): ν steps cost $\nu - 1$ applications of $\tilde{A}$, and the polynomial is fixed, so the smoother is a linear map. The runs use $\nu = 8$, $\kappa_s = 10$ with stout links ($M_5 = 1$) and $\nu = 12$, $\kappa_s = 30$ without ($M_5 = 1.8$).

4.2 The factorized coarse correction

The inverse of the right form. Invert the product one factor at a time:

$$\tilde{A}_R^{-1} = \left[(S_{\rm PV}C_{{\rm PV},e}^{-1})^\dagger (SC_e^{-1})\right]^{-1} \tag{4.17}$$

$$= (SC_e^{-1})^{-1}\left[(S_{\rm PV}C_{{\rm PV},e}^{-1})^\dagger\right]^{-1} \tag{4.18}$$

$$= C_eS^{-1}\left[(C_{{\rm PV},e}^{-1})^\dagger S_{\rm PV}^\dagger\right]^{-1} \tag{4.19}$$

$$= C_eS^{-1}\,(S_{\rm PV}^\dagger)^{-1}C_{{\rm PV},e}^\dagger \tag{4.20}$$

$$= C_e\;\pi D^{-1}\iota\;\; \pi(D_{\rm PV}^\dagger)^{-1}\iota\;\; C_{{\rm PV},e}^\dagger, \tag{4.21}$$

where Eq. (2.9) was used twice. The inverse is a product of the inverses of two *one-hop* operators on the whole lattice, D and $D_{\rm PV}^\dagger$, and site-local factors.

Galerkin replacement. With a prolongation $\mathcal{P}$ (coarse to whole lattice), replace each inverse by its Galerkin approximation:

$$X^{-1} \approx \mathcal{P}\big(\mathcal{P}^\dagger X\mathcal{P}\big)^{-1}\mathcal{P}^\dagger. \tag{4.22}$$

Both coarse operators $\mathcal{D}_{\rm c} = \mathcal{P}^\dagger D\mathcal{P}$ and $\mathcal{P}^\dagger D_{\rm PV}^\dagger\mathcal{P}$ are nearest-neighbour on the coarse lattice, because D is one hop and every aggregate contains the whole fifth dimension. The correction of a residual r (Eq. (4.21) read from the right) is

$$e = O_L\,\pi\mathcal{P}\,\mathcal{D}_{\rm c}^{-1}\mathcal{P}^\dagger\iota\; O_M\;\pi\mathcal{P}\big(\mathcal{P}^\dagger D_{\rm PV}^\dagger\mathcal{P}\big)^{-1}\mathcal{P}^\dagger\iota\; O_R\, r \tag{4.23}$$

with the site-local factors

form	O_R	O_M	O_L
right, any α	$C_{{\rm PV},e}^\dagger$	1	C_e
left, α in D only	1	$C_eQ_eC_{{\rm PV},e}^\dagger$	Q_e^{-1}
left, α in D and $D_{\rm PV}$	Q_e^{-1}	$C_eQ_e^2C_{{\rm PV},e}^\dagger$	Q_e^{-1}

In words: (i) put $O_R r$ on the whole lattice, restrict, solve with the adjoint Pauli–Villars coarse operator, prolong; (ii) take the even part, multiply by O_M, put it back on the whole lattice, restrict, solve with the quark coarse operator, prolong; (iii) take the even part and multiply by O_L. The coarse operators are always those of the plain D and D_{PV}, so one setup serves both forms and every α.

The left form with α, step by step. For α in D only,

$$\tilde{A}_L^{(\alpha)\,-1} = \left[(C_{\mathrm{PV},e}^{-1} S_{\mathrm{PV}})^\dagger\, Q_e^{-1} C_e^{-1} S Q_e\right]^{-1} \tag{4.24}$$

$$= \left(Q_e^{-1} C_e^{-1} S Q_e\right)^{-1} \left[(C_{\mathrm{PV},e}^{-1} S_{\mathrm{PV}})^\dagger\right]^{-1} \tag{4.25}$$

$$= Q_e^{-1} S^{-1} C_e Q_e \left[S_{\mathrm{PV}}^\dagger (C_{\mathrm{PV},e}^{-1})^\dagger\right]^{-1} \tag{4.26}$$

$$= Q_e^{-1} S^{-1} C_e Q_e\, C_{\mathrm{PV},e}^\dagger\, (S_{\mathrm{PV}}^\dagger)^{-1} \tag{4.27}$$

$$= \underbrace{Q_e^{-1}}_{O_L}\ \pi D^{-1} \iota\ \underbrace{C_e Q_e C_{\mathrm{PV},e}^\dagger}_{O_M}\ \pi (D_{\mathrm{PV}}^\dagger)^{-1} \iota\ \underbrace{1}_{O_R}. \tag{4.28}$$

Unlike the right form, the Galerkin factor of $D_{\mathrm{PV}}^\dagger$ acts on the plain residual, and $C_{\mathrm{PV},e}^\dagger$ lands between the two coarse solves, in O_M.

4.3 Aggregation and the chirality split

The fine lattice is divided into 4^4 blocks, each with the whole fifth dimension; each block is one coarse site. Every test vector v is split by $\Gamma_5 = \gamma_5 R_5$, $(R_5\psi)_s = \psi_{L_s-1-s}$:

$$v_\pm = \tfrac{1}{2}(1 \pm \Gamma_5) v, \qquad \Gamma_5 v_\pm = \pm v_\pm. \tag{4.29}$$

With n_v test vectors a coarse site has $2n_v$ components, the $+$ parts first. Each column of $\mathcal{P}$ is then an eigenvector of Γ_5, so with $\Sigma = \mathrm{diag}(\mathbf{1}_{n_v}, -\mathbf{1}_{n_v})$ per coarse site

$$\Gamma_5 \mathcal{P} = \mathcal{P} \Sigma. \tag{4.30}$$

Half the storage. Write the condition $\Gamma_5 v_\pm = \pm v_\pm$ out slice by slice:

$$(\Gamma_5 v_\pm)_s = \gamma_5\, (R_5 v_\pm)_s = \gamma_5\, v_\pm(L_s - 1 - s) \tag{4.31}$$

$$\pm v_\pm(s) = \gamma_5\, v_\pm(L_s - 1 - s) \tag{4.32}$$

$$v_\pm(L_s - 1 - s) = \pm \gamma_5\, v_\pm(s). \tag{4.33}$$

So each part is fixed by its slices $s < L_s/2$; $\mathcal{P}$ is stored only there ("folded"), in half precision. Restriction and prolongation unfold on the fly: for the restriction of a fine field f,

$$\left(\mathcal{P}^\dagger f\right)_{X,k} = \sum_{x \in X} \sum_{s < L_s/2} \mathcal{P}^*_{x,s;k}\, \left(f_s \pm \gamma_5 f_{L_s-1-s}\right) \tag{4.34}$$

(the sign that of the chirality of column k), which is what `fold_restrict_tiled_kernel` computes.

4.4 Orthonormalization per block: CholeskyQR2

In each block and chirality the columns $w_1, \ldots, w_n$ ($n = 2\times$ the number of filtered vectors, see below) are made orthonormal. Collect them as the columns of W and write the QR decomposition with a positive diagonal of R:

$$W = QR \tag{4.35}$$

$$W^\dagger W = R^\dagger Q^\dagger Q R = R^\dagger R \tag{4.36}$$

$$G \equiv W^\dagger W = R^\dagger R \quad \text{(Cholesky)} \tag{4.37}$$

$$Q = W R^{-1}. \tag{4.38}$$

So one Gram matrix ($n \times n$), one Cholesky factor and one triangular solve give Q; the factorization is unique, so this Q equals what modified Gram–Schmidt produces. In floating point one pass loses orthogonality like the square of the condition number of W, so the pass is done twice (CholeskyQR2): the second G is close to the identity and repairs the first. Each pass reads W twice, against one sweep per pair of columns for Gram–Schmidt (800 sweeps for $n = 40$). Where a pivot of the first pass falls below 10^{-12} of its diagonal (nearly dependent columns), the code falls back to Gram–Schmidt for those blocks.

4.5 The basis $\{s, \hat{t}\}$ and why the correction is exact

The setup produces vectors y_j on the even sites (next section). Each is lifted to the whole lattice in two ways, one for each coarse solve:

$$s = (C_e^{-1} y,\ -C_o^{-1} D_{oe} C_e^{-1} y), \tag{4.39}$$

$$\hat{t} = (a,\ -(C_{\mathrm{PV},o}^{\dagger})^{-1} D_{\mathrm{PV},eo}^{\dagger} a), \qquad a = \pi D s. \tag{4.40}$$

Ds vanishes on the odd sites.

$$Ds = \begin{pmatrix} C_e & D_{eo} \\ D_{oe} & C_o \end{pmatrix} \begin{pmatrix} C_e^{-1} y \\ -C_o^{-1} D_{oe} C_e^{-1} y \end{pmatrix} \tag{4.41}$$

$$= \begin{pmatrix} C_e C_e^{-1} y - D_{eo} C_o^{-1} D_{oe} C_e^{-1} y \\ D_{oe} C_e^{-1} y - C_o C_o^{-1} D_{oe} C_e^{-1} y \end{pmatrix} \tag{4.42}$$

$$= \begin{pmatrix} (C_e - D_{eo} C_o^{-1} D_{oe}) C_e^{-1} y \\ 0 \end{pmatrix} \tag{4.43}$$

$$= \iota\, S C_e^{-1} y. \tag{4.44}$$

In the same way $D_{\mathrm{PV}}^{\dagger} \hat{t} = \iota S_{\mathrm{PV}}^{\dagger} a$.

Exactness of the right form. Let $s, \hat{t} \in \operatorname{range} \mathcal{P}$ and solve both coarse systems exactly. The Galerkin map $\mathcal{P}(\mathcal{P}^{\dagger} X \mathcal{P})^{-1} \mathcal{P}^{\dagger} X$ is the identity on $\operatorname{range} \mathcal{P}$. Follow $r = \tilde{A}_R y$ through the correction ($O_R = C_{\mathrm{PV},e}^{\dagger}$, $O_M = 1$, $O_L = C_e$):

$$O_R r = C_{\mathrm{PV},e}^{\dagger} (S_{\mathrm{PV}} C_{\mathrm{PV},e}^{-1})^{\dagger} S C_e^{-1} y \tag{4.45}$$

$$= C_{\mathrm{PV},e}^{\dagger} (C_{\mathrm{PV},e}^{-1})^{\dagger} S_{\mathrm{PV}}^{\dagger} S C_e^{-1} y \tag{4.46}$$

$$= S_{\mathrm{PV}}^{\dagger}\, a, \qquad (a = S C_e^{-1} y = \pi D s) \tag{4.47}$$

$$\iota O_R r = \iota S_{\mathrm{PV}}^{\dagger} a = D_{\mathrm{PV}}^{\dagger} \hat{t} \tag{4.48}$$

$$\mathcal{P}\big(\mathcal{P}^{\dagger} D_{\mathrm{PV}}^{\dagger} \mathcal{P}\big)^{-1} \mathcal{P}^{\dagger}\, D_{\mathrm{PV}}^{\dagger} \hat{t} = \hat{t} \qquad (\hat{t} \in \operatorname{range} \mathcal{P}) \tag{4.49}$$

$$\pi \hat{t} = a = \pi D s \tag{4.50}$$

$$\iota O_M \pi \hat{t} = \iota \pi D s = D s \qquad (Ds \text{ is zero on the odd sites}) \tag{4.51}$$

$$\mathcal{P} \mathcal{D}_{\mathrm{c}}^{-1} \mathcal{P}^{\dagger} D s = s \qquad (s \in \operatorname{range} \mathcal{P}) \tag{4.52}$$

$$e = O_L \pi s = C_e C_e^{-1} y = y. \tag{4.53}$$

So the correction inverts $\tilde{A}_R$ exactly on the span of the y_j. In practice $\mathcal{P}$ is in half precision and the coarse solves are inexact (relative residual 10^{-1}), and the left form is not exact at all; the cycle converges nonetheless.

4.6 The coarse operators by probing

$\mathcal{D}_c = \mathcal{P}^\dagger D \mathcal{P}$ has on each coarse site X a $2n_v \times 2n_v$ diagonal block $A(X)$ and one link block $B_\mu^\pm(X)$ per direction: nine blocks. They are found column by column. For component k, put a unit vector ϵ_k on *every* coarse site and prolong: $\phi = \mathcal{P}\epsilon_k$. Split $D = C + H$ (H the hops) and H by leg (direction and sign), and split each leg into the part that stays inside its block and the part that crosses to the neighbour block:

$$D\phi = C\phi + \sum_{\text{legs } l} H_l \phi \tag{4.54}$$

$$= \underbrace{C\phi + \sum_l H_l^{\text{in}}\phi}_{\text{stays in } X} + \sum_l \underbrace{H_l^{\text{cross}}\phi}_{\text{comes from } X\pm\hat{\mu}} \quad . \tag{4.55}$$

Restricting, $\left(\mathcal{P}^\dagger(C\phi + \sum_l H_l^{\text{in}}\phi)\right)_X$ is column k of $A(X)$, and $\left(\mathcal{P}^\dagger H_l^{\text{cross}}\phi\right)_X$ is column k of the link block of leg l, since every coarse site carries the same unit vector and a crossing hop brings exactly the neighbour's contribution. Doing this for all $2n_v$ components fills all nine blocks (`coarse_blocks`, `hop_split`, `restrict_leg`). Several components are probed at a time, as many as a quarter of the free device memory allows (four on 24^4), so that $\mathcal{P}$ is read once for them.

4.7 The coarse solves

Quark factor. $\mathcal{D}_c c = g$ by BiCGStab on its even-odd Schur complement, preconditioned by the inverse diagonal blocks. With $\mathcal{D}_c = \begin{pmatrix} A_{ee} & B_{eo} \\ B_{oe} & A_{oo} \end{pmatrix}$ on the coarse lattice, exactly as in Eq. (2.7),

$$(A_{ee} - B_{eo}A_{oo}^{-1}B_{oe})c_e = g_e - B_{eo}A_{oo}^{-1}g_o, \tag{4.56}$$

$$c_o = A_{oo}^{-1}(g_o - B_{oe}c_e). \tag{4.57}$$

Pauli–Villars factor. $(\mathcal{P}^\dagger D_{\text{PV}}^\dagger \mathcal{P})w = g$. With $\Gamma_5\mathcal{P} = \mathcal{P}\Sigma$,

$$\mathcal{P}^\dagger(\Gamma_5 D_{\text{PV}}\Gamma_5)\mathcal{P} = \mathcal{P}^\dagger\Gamma_5\, D_{\text{PV}}\, \Gamma_5\mathcal{P} \tag{4.58}$$

$$= (\Gamma_5\mathcal{P})^\dagger\, D_{\text{PV}}\, (\Gamma_5\mathcal{P}) \tag{4.59}$$

$$= \Sigma\mathcal{P}^\dagger D_{\text{PV}}\mathcal{P}\Sigma = \Sigma\mathcal{D}_c^{\text{PV}}\Sigma. \tag{4.60}$$

If $\Gamma_5 D_{\text{PV}}\Gamma_5 = D_{\text{PV}}^\dagger$ (true for $c = 0$, off by $O(c)$ otherwise, Sec. 6), the adjoint coarse factor is $\Sigma\mathcal{D}_c^{\text{PV}}\Sigma$, and

$$\Sigma\mathcal{D}_c^{\text{PV}}\Sigma\, w = g \tag{4.61}$$

$$\mathcal{D}_c^{\text{PV}}\Sigma\, w = \Sigma g \qquad (\Sigma^2 = 1) \tag{4.62}$$

$$\mathcal{H}_c\, w = \Sigma g, \qquad \mathcal{H}_c \equiv \mathcal{D}_c^{\text{PV}}\Sigma, \tag{4.63}$$

$\mathcal{H}_c$ Hermitian (indefinite) for $c = 0$, solved by MINRES with the absolute value of the diagonal blocks as a positive preconditioner. Both factors stop at $|r|^2/|b|^2 < 10^{-2}$.

Levels below. At small masses the coarse solves themselves become expensive (on 24^4 at $m = 0.001$ about 2000 applications of $\mathcal{D}_c$ per correction). The coarse sites are then grouped into super-blocks (2, then 3), local vectors are formed from $\mathcal{P}^\dagger s_j$, $\mathcal{P}^\dagger \hat{t}_j$, and a V-cycle preconditions the coarse Krylov solve: with Q_ℓ the prolongation into level $\ell - 1$ and B_ℓ its diagonal blocks,

$$z = Q_{\ell+1}M_{\ell+1}Q_{\ell+1}^\dagger r, \qquad M_\ell r = z + B_\ell^{-1}(r - \mathcal{D}_\ell z), \tag{4.64}$$

the last level inverted densely. This brings the coarse cost down to about 50 applications per correction.

5 Setup

The setup produces the vectors y_j, then $\{s_j, \hat{t}_j\}$, $\mathcal{P}$, the coarse operators and the levels. It always uses the right form.

5.1 Chebyshev-filtered subspace iteration

The y_j should approximate the low right singular vectors of $\tilde{A}_R$, i.e. the low eigenvectors of $H_n = \tilde{A}_R^\dagger \tilde{A}_R$. With $\ell(x) = (2x - (\lambda_{\rm hi} + \lambda_{\rm lo}))/(\lambda_{\rm hi} - \lambda_{\rm lo})$, the polynomial $T_d(\ell(x))$ is bounded by 1 on $[\lambda_{\rm lo}, \lambda_{\rm hi}]$ and grows fast below $\lambda_{\rm lo}$. One round on a block X of m vectors:

$$X \leftarrow T_d(\ell(H_n))X \qquad \text{(three-term recurrence, } d \text{ applications of } H_n) \tag{5.1}$$

$$X \leftarrow \text{orthonormalize } X \tag{5.2}$$

$$G = X^\dagger H_n X = V \Theta V^\dagger, \quad X \leftarrow XV \qquad \text{(Rayleigh–Ritz)} \tag{5.3}$$

$$\lambda_{\rm lo} \leftarrow \theta_{\max}. \tag{5.4}$$

The degree is $d = 1.5\sqrt{\lambda_{\rm hi}/\lambda_{\rm lo}}$, between 40 and 1000; the first round uses $\lambda_{\rm lo} = \lambda_{\rm hi}/1000$. *Observation* ($24^4$, $m = 0.001$, 20 vectors): after the first round (degree 48) the largest Ritz value is still $0.34\,\lambda_{\rm hi}$, so the second round has degree 40 and damps only the top two thirds of the spectrum. The filter is weak; the bootstrap does most of the work.

5.2 Bootstrap

Each round: build $\mathcal{P}$ and the cycle $\mathcal{M}$ from the current y; replace every y_j by four steps of FGMRES on $\tilde{A}_R x = y_j$ from $x = 0$ preconditioned by $\mathcal{M}$; orthonormalize y through its Gram matrix. Why it works: write y_j in eigenvectors of the preconditioned operator. The cycle already handles the components inside its coarse space, and FGMRES with it removes them first; what remains, and is therefore amplified relative to the rest, is what the coarse space misses. In HMC the levels below are used from the second round on (in the first, on the filtered vectors, the coarse solves were as fast without them).

5.3 Cost and memory

On 24^4 ($L_s = 8$, 20 filtered vectors, $m = 0.001$, A30): the filter about 35 s and the bootstrap about 77 s, the setup in all about 113 s. A bootstrap round makes its correction ($\mathcal{P}$ 1.8 s, the two coarse operators about 4.4 s each, the inverse blocks 0.8 s, the top of the smoother 1.3 s; the levels 7 s more in the second round) and then improves the 20 vectors (about 20 s in the first round, 26 s in the second). $\mathcal{P}$ (40 columns, half precision, whole lattice) is the largest object, about 4.7 GB.

6 Solving with $D^\dagger$

HMC needs $(D^\dagger D)^{-1}$, so it solves with $D^\dagger$ as often as with D. A second setup for $D^\dagger$ would double the cost. The tool that avoids it is an exact relation between D and $D^\dagger$.

6.1 $\Gamma_5 D \Gamma_5 \neq D^\dagger$

Each building block transforms simply under $\Gamma_5 = \gamma_5 R_5$:

$$\Gamma_5 W \Gamma_5 = W^\dagger, \qquad \Gamma_5 D_\pm \Gamma_5 = D_\pm^\dagger, \qquad \Gamma_5 Y \Gamma_5 = Y^\dagger. \tag{6.1}$$

Then, from $D = D_+ - D_-Y$,

$$\Gamma_5 D \Gamma_5 = \Gamma_5 D_+ \Gamma_5 - \Gamma_5 D_- \Gamma_5 \, \Gamma_5 Y \Gamma_5 = D_+^\dagger - D_-^\dagger Y^\dagger, \tag{6.2}$$

$$D^\dagger = D_+^\dagger - (D_- Y)^\dagger = D_+^\dagger - Y^\dagger D_-^\dagger. \tag{6.3}$$

The order of the last product differs. Subtracting,

$$\Gamma_5 D \Gamma_5 - D^\dagger = Y^\dagger D_-^\dagger - D_-^\dagger Y^\dagger \tag{6.4}$$

$$= Y^\dagger (cW^\dagger - 1) - (cW^\dagger - 1) Y^\dagger \tag{6.5}$$

$$= c\,(Y^\dagger W^\dagger - W^\dagger Y^\dagger) = c\,[Y^\dagger, W^\dagger], \tag{6.6}$$

which is of order c (zero for Shamir, 0.15 relative on a 12^4 stout configuration at $c = 0.5$).

6.2 The exact relation

D_+ and D_- are functions of W, so they commute. Factor D_- out to the left (formally):

$$D = D_- K_D, \qquad K_D = D_-^{-1} D_+ - Y. \tag{6.7}$$

K_D is Γ_5-Hermitian:

$$\Gamma_5 K_D \Gamma_5 = \Gamma_5 D_-^{-1} \Gamma_5 \, \Gamma_5 D_+ \Gamma_5 - \Gamma_5 Y \Gamma_5 \tag{6.8}$$

$$= (D_-^\dagger)^{-1} D_+^\dagger - Y^\dagger \tag{6.9}$$

$$= D_+^\dagger (D_-^\dagger)^{-1} - Y^\dagger \qquad (D_\pm^\dagger \text{ commute}) \tag{6.10}$$

$$= (D_-^{-1} D_+)^\dagger - Y^\dagger = K_D^\dagger. \tag{6.11}$$

So

$$D^\dagger = K_D^\dagger D_-^\dagger \tag{6.12}$$

$$= \Gamma_5 K_D \Gamma_5 \, D_-^\dagger \tag{6.13}$$

$$= \Gamma_5 D_-^{-1} D \, \Gamma_5 \, D_-^\dagger, \tag{6.14}$$

and multiplying by $D_- \Gamma_5$ from the left,

$$\boxed{D_- \, \Gamma_5 \, D^\dagger = D \, \Gamma_5 \, D_-^\dagger} \tag{6.15}$$

with no inverse left. D_- itself cannot be inverted in practice: $D_- = c(W - 1/c)$ is c times a Wilson operator of mass $-M_5 - 1/c = -3$, whose spectrum surrounds the origin (GCR stalls at $|r|^2/|b|^2 \approx 0.74$).

6.3 $\Gamma_5 D^{-1} \Gamma_5$ as a preconditioner of $D^\dagger$

Multiply $D^\dagger$ from the right by $\Gamma_5 D^{-1} \Gamma_5$ and use the relation above in the form $D^\dagger = \Gamma_5 D_-^{-1} D \, \Gamma_5 D_-^\dagger$, together with $\Gamma_5 D_-^\dagger \Gamma_5 = D_-$:

$$D^\dagger \, \Gamma_5 D^{-1} \Gamma_5 = \Gamma_5 D_-^{-1} D \, \Gamma_5 D_-^\dagger \, \Gamma_5 \, D^{-1} \Gamma_5 \tag{6.16}$$

$$= \Gamma_5 D_-^{-1} D \, D_- \, D^{-1} \, \Gamma_5 \tag{6.17}$$

$$= \Gamma_5 D_-^{-1} (D_- D + [D, D_-]) D^{-1} \Gamma_5 \tag{6.18}$$

$$= \Gamma_5 (1 + D_-^{-1} [D, D_-] D^{-1}) \Gamma_5. \tag{6.19}$$

The commutator, with $D = D_+ - D_- Y$ and $[D_+, D_-] = 0$:

$$[D, D_-] = [D_+, D_-] - [D_- Y, D_-] \tag{6.20}$$
$$= -(D_- Y D_- - D_- D_- Y) = -D_- [Y, D_-] \tag{6.21}$$
$$= -D_- [Y, cW - 1] = -c\, D_- [Y, W]. \tag{6.22}$$

The factor D_- cancels against D_-^{-1}:

$$\boxed{D^\dagger \Gamma_5 D^{-1} \Gamma_5 = \Gamma_5 (1 - N) \Gamma_5, \qquad N = c\,[Y, W]\, D^{-1}} \tag{6.23}$$

$[Y, W] = -\frac{1}{2}[Y, K]$ is bounded (one hop in four dimensions, one in the fifth), so N is large only on the low modes of D. There the commutator flips chirality and moves the mode to the other wall, out of the low-mode span, so N^2 is small and $(1 - N)^{-1} \approx 1 + N$: GCR needs only a few more steps than for the identity. Measured: 2–3 iterations more than the solve with D.

6.4 On the even sites, step by step

The solves of HMC are on the even sites. Since Γ_5 acts within a site it commutes with π and ι, and with Eq. (2.9)

$$S^{-\dagger} = \pi D^{-\dagger} \iota \approx \pi \Gamma_5 D^{-1} \Gamma_5 \iota = \Gamma_5\, \pi D^{-1} \iota\, \Gamma_5 = \Gamma_5 S^{-1} \Gamma_5. \tag{6.24}$$

The solve of $D^\dagger w = r$:

$$r'_e = r_e - D^\dagger_{oe} (C^\dagger_o)^{-1} r_o \qquad \text{so that } S^\dagger w_e = r'_e, \tag{6.25}$$
$$w_e = (C^\dagger_e)^{-1} z \qquad \Longrightarrow\ S^\dagger (C^\dagger_e)^{-1} z = r'_e \iff K^\dagger z = r'_e, \tag{6.26}$$
$$z = \text{GCR on } K^\dagger \text{ with preconditioner } C^\dagger_e\, \Gamma_5 S^{-1} \Gamma_5, \tag{6.27}$$
$$S^{-1} t \approx C_e^{-1}\, \mathcal{M}\, (S_{\rm PV} C^{-1}_{{\rm PV},e})^\dagger t \qquad \text{(one cycle of } \tilde{A}_R\text{)}, \tag{6.28}$$
$$w_e = (C^\dagger_e)^{-1} z, \qquad w_o = (C^\dagger_o)^{-1} (r_o - D^\dagger_{eo} w_e). \tag{6.29}$$

The line $K^\dagger z = r'_e$ uses $K = C_e^{-1} S$: $K^\dagger = S^\dagger (C^\dagger_e)^{-1}$. The approximation for S^{-1} comes from Eq. (4.21): $S^{-1} = C_e^{-1} \tilde{A}_R^{-1} (S_{\rm PV} C^{-1}_{{\rm PV},e})^\dagger$. In the preconditioner, $(S_{\rm PV} C^{-1}_{{\rm PV},e})^\dagger t$ is the right-hand side of $\tilde{A}_R$ for the source $(t, 0)$ (no odd part), so it is applied directly as one adjoint Schur factor of $D_{\rm PV}$. Each GCR step costs one cycle, as for D.

7 HMC

7.1 The even-odd Ratio action

Two flavours over two Pauli–Villars flavours, as Bridge++ writes it (`Action_F_alt_Ratio_eo`): with $K = C_e^{-1} S$ (the left Schur factor) and $H = \Gamma_5 K$,

$$S_F = \phi^\dagger_e H_{\rm PV}\, (H^\dagger H)^{-1} H^\dagger_{\rm PV} \phi_e. \tag{7.1}$$

Since Γ_5 is Hermitian and unitary,

$$H^\dagger H = K^\dagger \Gamma_5 \Gamma_5 K = K^\dagger K, \tag{7.2}$$

and with $v_1 \equiv H^\dagger_{\rm PV} \phi_e = K^\dagger_{\rm PV} \Gamma_5 \phi_e$,

$$S_F = v_1^\dagger\, (K^\dagger K)^{-1}\, v_1. \tag{7.3}$$

Heat bath. Draw ξ_e Gaussian and set $\phi_e = H_{\mathrm{PV}}(H_{\mathrm{PV}}^\dagger H_{\mathrm{PV}})^{-1}H^\dagger\xi_e$. Then, step by step,

$$H_{\mathrm{PV}}^\dagger\phi_e = H_{\mathrm{PV}}^\dagger H_{\mathrm{PV}}(H_{\mathrm{PV}}^\dagger H_{\mathrm{PV}})^{-1}H^\dagger\xi_e = H^\dagger\xi_e, \tag{7.4}$$

$$S_F = (H^\dagger\xi_e)^\dagger(H^\dagger H)^{-1}(H^\dagger\xi_e) \tag{7.5}$$

$$= \xi_e^\dagger H(H^\dagger H)^{-1}H^\dagger\xi_e \tag{7.6}$$

$$= \xi_e^\dagger HH^{-1}H^{-\dagger}H^\dagger\xi_e = |\xi_e|^2, \tag{7.7}$$

as required. The heat bath costs one solve with the Pauli–Villars operator (CG; it is well conditioned).

Force. Vary the links. With $\chi = \Gamma_5\phi_e$, $v_1 = K_{\mathrm{PV}}^\dagger\chi$ and $v_2 = (K^\dagger K)^{-1}v_1$:

$$\delta S_F = \delta v_1^\dagger\, v_2 + v_2^\dagger\,\delta v_1 + v_1^\dagger\,\delta[(K^\dagger K)^{-1}]v_1 \tag{7.8}$$

$$= 2\,\Re\,(\delta v_1)^\dagger v_2 - v_1^\dagger(K^\dagger K)^{-1}\,\delta(K^\dagger K)\,(K^\dagger K)^{-1}v_1 \tag{7.9}$$

$$= 2\,\Re\,\chi^\dagger\,\delta K_{\mathrm{PV}}\,v_2 - v_2^\dagger(\delta K^\dagger\,K + K^\dagger\,\delta K)v_2 \tag{7.10}$$

$$= 2\,\Re\,\chi^\dagger\,\delta K_{\mathrm{PV}}\,v_2 - 2\,\Re\,(Kv_2)^\dagger\,\delta K\,v_2. \tag{7.11}$$

So every force evaluation needs one solve $v_2 = (K^\dagger K)^{-1}v_1$, and every energy evaluation (at the start and the end of a trajectory) one more. The derivative of K is that of the hops in S (C_e does not depend on the links), taken with respect to the smeared links and carried back through the stout derivative.

7.2 $(K^\dagger K)^{-1}$ with the multigrid: two solves

Invert the product and use $K = C_e^{-1}S$:

$$(K^\dagger K)^{-1} = K^{-1}K^{-\dagger} \tag{7.12}$$

$$= (C_e^{-1}S)^{-1}\big[(C_e^{-1}S)^\dagger\big]^{-1} \tag{7.13}$$

$$= S^{-1}C_e\,C_e^\dagger S^{-\dagger}. \tag{7.14}$$

So $z = (K^\dagger K)^{-1}r$ is computed as

$$w = C_e^\dagger\,S^{-\dagger}r = C_e^\dagger\,\pi D^{-\dagger}\iota\,r \qquad \text{(a solve with } D^\dagger\text{, Sec. 6)}, \tag{7.15}$$

$$z = S^{-1}C_e\,w = \pi D^{-1}\iota\,C_e w \qquad \text{(a solve with } D\text{)}. \tag{7.16}$$

Each solve stops at a relative residual 10^{-10}. The normal equation must be solved much more accurately (10^{-16} for the force, 10^{-24} for the energy), so the two solves are repeated on the residual (iterative refinement):

$$r^{(k)} = b - K^\dagger K z^{(k)} \tag{7.17}$$

$$z^{(k+1)} = z^{(k)} + K^{-1}K^{-\dagger}r^{(k)} \qquad \text{(two solves to } 10^{-10}\text{)}, \tag{7.18}$$

each sweep gaining about ten digits, at most eleven sweeps.

When the sweeps stall. If the sweeps stop gaining (the coarse space no longer fits the links, as after large link changes), the solve is repeated from the current z with $\mathcal{P}$ and the coarse operators rebuilt for the current links, then also after a bootstrap; if it still stalls, CG on $K^\dagger K$ finishes it from the current z (with at least 50000 iterations), with a warning. A trajectory is never lost to the multigrid.

7.3 Neff's α in HMC

CG. From Eq. (2.20), $K = Q_e K^{(\alpha)} Q_e^{-1}$. Then

$$K^\dagger K = Q_e^{-1} K^{(\alpha)\dagger} Q_e Q_e K^{(\alpha)} Q_e^{-1} \tag{7.19}$$

$$= Q_e^{-1} (Q_e K^{(\alpha)})^\dagger (Q_e K^{(\alpha)}) Q_e^{-1} \tag{7.20}$$

$$= Q_e^{-1} \, (KQ_e)^\dagger (KQ_e) \, Q_e^{-1} \qquad (KQ_e = Q_e K^{(\alpha)}), \tag{7.21}$$

$$(K^\dagger K)^{-1} = Q_e \big[(KQ_e)^\dagger (KQ_e) \big]^{-1} Q_e. \tag{7.22}$$

So CG runs on the normal equation of KQ_e, whose singular values are close to those of $\alpha K^{(\alpha)}$, and the action and the force stay those of K. On 24^4 at $m = 0.001$ this cut the CG iterations by 22%.

Multigrid. The solve with D uses $D^{-1} = Q D^{(\alpha)\,-1}$: the left form $\tilde{A}_L^{(\alpha)}$ with the factors of the table above and $\psi = Q\psi^{(\alpha)}$. The solve with $D^\dagger$ is unchanged: $D^{(\alpha)\dagger} = QD^\dagger$, so α would only rescale its right-hand side. The setup is that of D.

Why α helps CG more than the multigrid. $K^{(\alpha)} = Q_e^{-1} K Q_e$ has the eigenvalues of K but other singular values (Q_e is not unitary). CG on the normal equation sees the singular values, so it gains. The multigrid's iterations do not follow the whole condition number: the coarse grid takes the low modes and the smoother the rest, and α helps only the smoother's part. In the propagator mass scan this saves 2–3 iterations at every mass with $M_5 = 1.8$ (smoother interval $[\lambda_{\text{hi}}/30, \lambda_{\text{hi}}]$, wide) and only at heavy masses with stout links ($[\lambda_{\text{hi}}/10, \lambda_{\text{hi}}]$).

7.4 What is kept along a trajectory

- The test vectors y: kept from solve to solve (on the host between solves), improved by a bootstrap (from the old y, not a new filter) when the pseudofermions are drawn.
- $\mathcal{P}$, the coarse operators, the top of the smoother's interval: kept from solve to solve while the links move, and rebuilt for the current links when the GCR iterations per solve have grown by half (`rebuild_iteration_growth` = 1.5) or after a bootstrap.
- The cycle itself is made for each solve from the kept parts and dropped after it.

7.5 Memory (24^4, $m = 0.001$, 20 vectors, levels below, A30)

part	GB
HMC itself (links, smeared levels, momenta, pseudofermion, operators)	3.3
multigrid kept between solves ($\mathcal{P}$ 4.75)	6.6
fields of the sweeps	2.6
the solve with $D^\dagger$ on top (GCR, Schur factors, cycle)	6.4
peak (live)	19.1

7.6 What the measurements say so far

- On a thermalized dynamical 16^4 configuration the multigrid reproduces the trajectories of CG (plaquette to 10^{-11}).
- There CG needs about 300 iterations per solve at $m = 0.01$ and at $m = 0.001$ alike: the low end of the spectrum is held up by the finite volume, not by the mass. The multigrid needs about 55 iterations, but one of them (a whole cycle) costs about 24 CG iterations, so it is 3–4 times slower per trajectory. Tuning gains at most 6%.

- On the quenched 24^4 configuration CG needs about 6400 iterations and the multigrid wins, but a trajectory of length 1 there blows up. So there are no data yet that show the multigrid faster in HMC; it needs a volume where CG takes more than about 1300 iterations per solve.

8 Map to the code

step	where
C, C^{-1}, hops, Γ_5, Q on the device	`src/backend/cuda_bridge/mdwf.rs`; `cuda_bridge/include/fermion/`
even-odd forms, $\tilde{A}_R$, $\tilde{A}_L$ (`PvRatio`, `SchurFactor`, `AfterQ`)	`src/fermion/even_odd.rs`
GCR, FGMRES, BiCGStab, MINRES, CG	`src/solver/`
Chebyshev smoother, two-grid cycle	`src/solver/preconditioner/multigrid/cycle.rs`; `src/backend/cuda_bridge/multigrid.rs`
factorized correction	`src/solver/preconditioner/multigrid/correction.rs`; `src/backend/cuda_bridge/correction.rs`
$\mathcal{P}$: fold, restrict, prolong, CholeskyQR2	`src/backend/cuda_bridge/coarse.rs`; `cuda_bridge/include/multigrid/fold.cuh`
coarse operators by probing	`CudaMdwf::coarse_blocks` in `src/backend/cuda_bridge/mdwf.rs`
levels below	`src/backend/cuda_bridge/levels.rs`
filter, bootstrap	`low_modes` in `src/solver/preconditioner/multigrid/mdwf_even_odd.rs`; `bootstrap` in `src/backend/cuda_bridge/multigrid.rs`
solve with $D^\dagger$	`adjoint_by_gamma5` in `src/fermion/even_odd.rs`
HMC: two solves, sweeps, retries	`src/hmc/dwf_multigrid.rs`; `src/backend/cuda_bridge/hmc_multigrid.rs`
HMC: action, force, CG with α, CG fallback	`src/hmc/even_odd.rs`; `src/hmc/solve.rs`